

QEMU & LINUX Support for RV32 Compat Mode

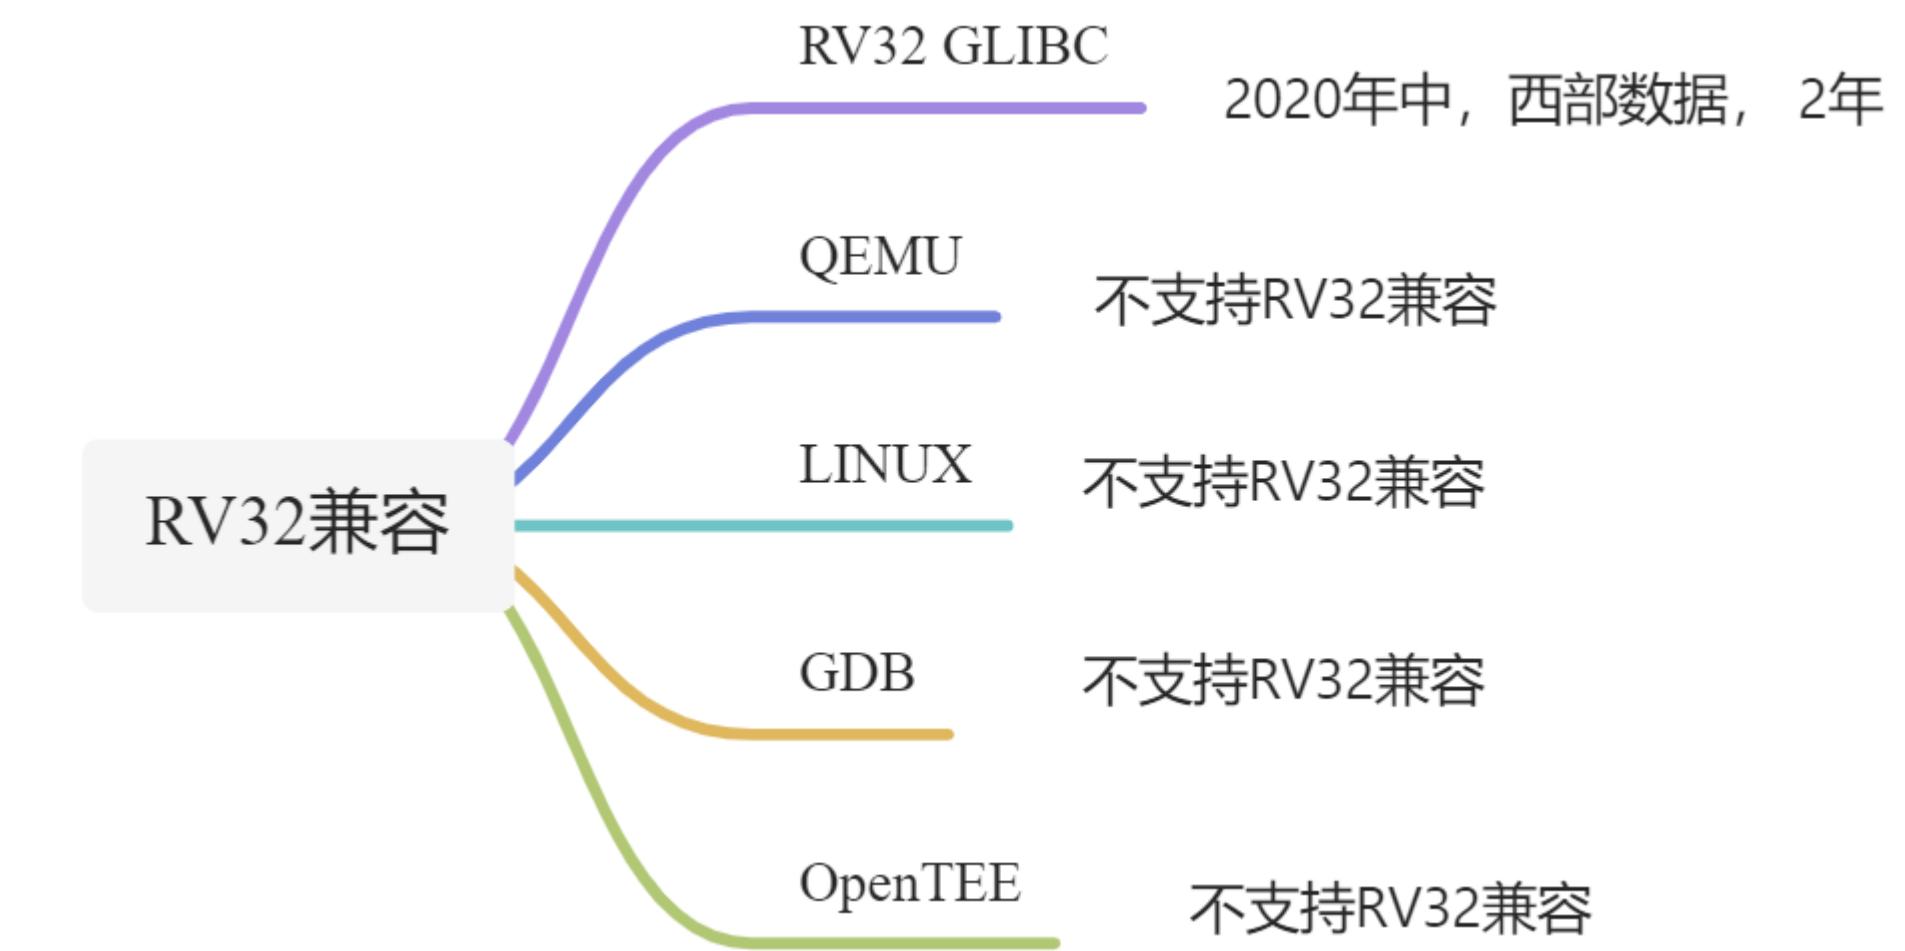
Liu Zhiwei & Guo Ren
T-Head Software Engineer
2022/08/24

Question

- Can we run RV32 application on 64-bit RISC-V CPU
 - Interesting & Complicated
 - Most architectures can run 32-bit application on 64-bit CPU
 - RISC-V does not have so many legacy 32-bit application

Exploring

- Prerequisites
 - Compiler for RV32 is ready
- Still need to exploring
 - Understanding the specification
 - Build an executive environment on QEMU
 - Necessary Linux kernel support



Run RV32 rootfs on RV64 Linux kernel

```

→ build ./qemu-system-riscv64 -cpu rv64 -D 1.log -M virt -m 1024m -nographic -bios ../fw_jump.bin -kernel
root=/dev/vda ro console=ttyS0 earlycon=sbi norandmaps" -netdev user,id=net0 -device virtio-net-device,netc
...
OpenSBI v1.0-34-g7924a0b220b7
...
[    0.000000] Linux version 5.19.0+ (guoren@fedora) (riscv64-unknown-linux-gnu-gcc (g5964b5cd7272) 11.1.0,
[    0.000000] OF: fdt: Ignoring memory range 0x40000000 - 0x40200000
[    0.000000] Machine model: riscv-virtio_qemu
[    0.000000] earlycon: sbio at I/O port 0x0 (options '')
[    0.000000] printk: bootconsole [sbio] enabled
[    0.000000] efi: UEFI not found.
[    0.000000] Zone ranges:
[    0.000000]   DMA32    [mem 0x0000000040200000-0x000000007fffffff]
[    0.000000]   Normal    empty
[    0.000000] Movable zone start for each node
[    0.000000] Early memory node ranges
[    0.000000]   node  0: [mem 0x0000000040200000-0x000000007fffffff]
[    0.000000]   Initmem setup node 0 [mem 0x0000000040200000-0x000000007fffffff]
...
[    0.000000] Virtual kernel memory layout:
[    0.000000]   fixmap : 0xff1bfffffee00000 - 0xff1bfffff000000 (2048 kB)
[    0.000000]   pci io : 0xff1bfffff000000 - 0xff1c000000000000 ( 16 MB)
[    0.000000]   vmemmap : 0xff1c000000000000 - 0xff20000000000000 (1024 TB)
[    0.000000]   vmalloc : 0xff20000000000000 - 0xff60000000000000 (16384 TB)
[    0.000000]   lowmem : 0xff60000000000000 - 0xff600003fe00000 (1022 MB)
[    0.000000]   kernel : 0xfffffff80000000 - 0xffffffffffffffff (2047 MB)
[    0.000000] Memory: 1010044K/1046528K available (7985K kernel code, 4921K rwdata, 4096K rodata, 2171K in
...
[    0.777981] EXT4-fs (vda): warning: mounting unchecked fs, running e2fsck is recommended
[    0.785870] EXT4-fs (vda): re-mounted. Quota mode: disabled.
Starting syslogd: OK
Starting klogd: OK
Running sysctl: OK
Initializing random number generator: OK
Saving random seed: [    5.521212] random: crng init done
Starting network: udhcpc: started, v1.35.0
udhcpc: broadcasting discover
udhcpc: broadcasting select for 10.0.2.15, server 10.0.2.2
udhcpc: lease of 10.0.2.15 obtained from 10.0.2.2, lease time 86400
deleting routers
adding dns 10.0.2.3

Welcome to Buildroot
buildroot login: root
# uname -a
Linux buildroot 5.19.0+ #23 SMP Sun Aug 7 19:58:52 EDT 2022 riscv64 GNU/Linux
# ls /lib/
ld-linux-riscv32-ilp32d.so.1  libgcc_s.so.1
libanl.so.1                  libm.so.6
libatomic.so                 libnss_dns.so.2
libatomic.so.1                libnss_files.so.2
libatomic.so.1.2.0             libpthread.so.0
libc.so.6                     libresolv.so.2
libcrypt.so.1                librt.so.1
libdl.so.2                   libutil.so.1
libgcc_s.so                  modules
# cat /proc/1/maps
55554000-55630000 r-xp 00000000 fe:00 17                           /bin/busybox
55631000-55633000 r--p 000dc000 fe:00 17                           /bin/busybox
55633000-55634000 rw-p 000de000 fe:00 17                           /bin/busybox
55634000-55655000 rw-p 00000000 00:00 0                            [heap]
77e8e000-77fbe000 r-xp 00000000 fe:00 137                         /lib/libc.so.6
77fbe000-77fc0000 r--p 0012f000 fe:00 137                         /lib/libc.so.6
77fc0000-77fc1000 rw-p 00131000 fe:00 137                         /lib/libc.so.6
77fc1000-77fc2000 rw-p 00000000 00:00 0
77fc2000-77fd3000 r-xp 00000000 fe:00 146                         /lib/libresolv.so.2
77fd3000-77fd4000 ---p 00008000 fe:00 146                         /lib/libresolv.so.2
77fd4000-77fd5000 r--p 00008000 fe:00 146                         /lib/libresolv.so.2
77fd5000-77fd6000 rw-p 00009000 fe:00 146                         /lib/libresolv.so.2
77fd6000-77fd8000 rw-p 00000000 00:00 0
77fd8000-77fd9000 r--p 00000000 00:00 0
77fd9000-77fdb000 r-xp 00000000 00:00 0
77fdb000-77ffc000 r--p 00000000 fe:00 132                         /lib/ld-linux-riscv32-ilp32d.so.1
77ffc000-77ffe000 r--p 00020000 fe:00 132                         /lib/ld-linux-riscv32-ilp32d.so.1
77ffe000-77fff000 rw-p 00022000 fe:00 132                         /lib/ld-linux-riscv32-ilp32d.so.1
7ffde000-7ffff000 rw-p 00000000 00:00 0                           [stack]

```

Contents 目录

- 01 Dynamic XLEN
- 02 UXLEN support on QEMU
- 03 RV64 LINUX COMPAT for 32bit U-mode

Dynamic XLEN

01

Specification Rules

- When XLEN < the widest supported XLEN
 - Ignore source operand register bits above the configured XLEN.
 - Sign-extend results to fill the entire widest supported XLEN in the destination register.
 - PC bits above XLEN are ignored, and when the pc is written, it is sign-extended to fill the widest supported XLEN.
- CSR width modulation
 - Zero extension or truncate
 - Specially specified

Specification not clarified

- Not seamless

It's important to understand that the defaults were not written to facilitate actively changing the effective XLEN for software "on the go". Rather, the expectation is that whatever software was running before is likely "done", and something else will now be run, without much carryover in the CSRs between the two, if indeed any(From John Hauser)

- $\text{MXLEN} \geq \text{SXLEN} \geq \text{UXLEN}$

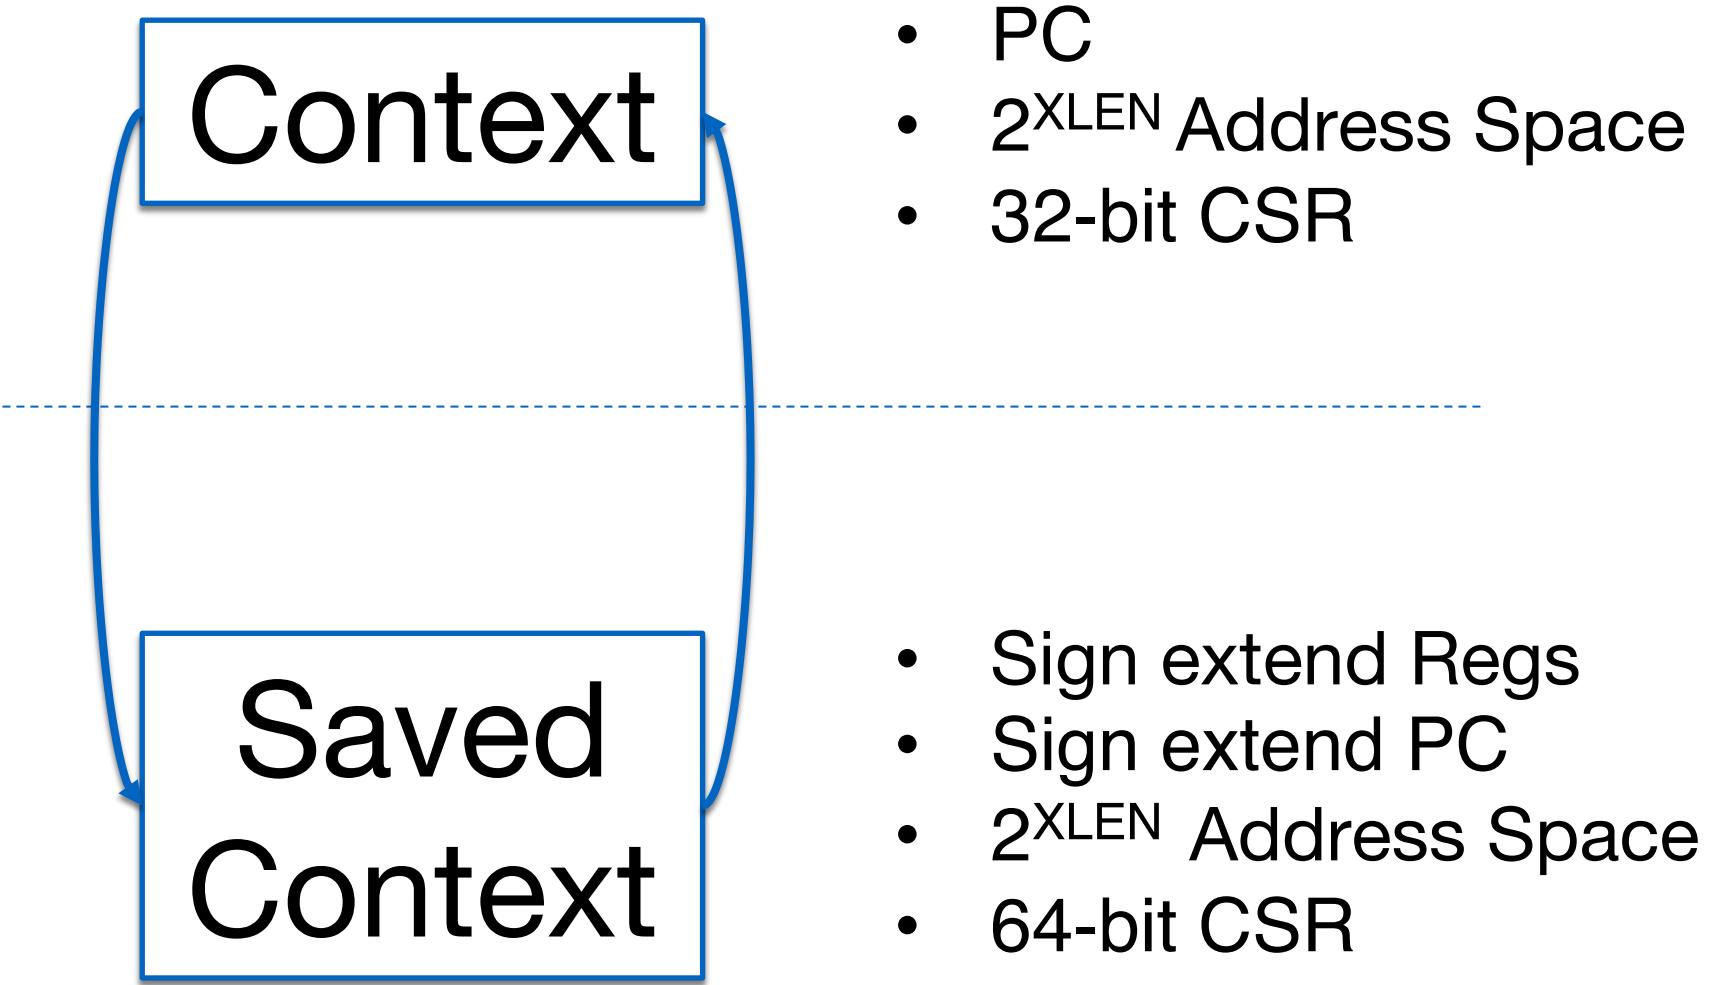
UXLEN support on QEMU

02

Ensure Context Switch Right

U-Mode(32bit)

S-Mode(64bit)



Implementation

CSR Modulation

- User mode CSR
- Split XLEN related fields

32-bit Instruction

- Use 64 bit TCG op to emulate 32-bit instruction
- Sign extend or Zero extend the source and always sign extend the destination
- Define OLEN to process the *.w

PC Calculation

- Always sign extend for branch.
- When exception, always sign extend the PC.
- When restore PC in runtime, always sign extend the PC.
- When fetch PC and decoding , always mask the high bits

Address Calculation

- Module 2^{XLEN}
- Always use before pointer masking

Upstream Process

UXLEN RFC
Patch

[\[RFC PATCH 00/13\] Support
UXL field in mstatus](#)



2021-08-09

Frédéric Pérot RV128

[\[PATCH 0/8\] RISC V partial
support for 128-bit
architecture](#)



2021-08-30

Richard Henderson MXL

[\[RFC PATCH 00/13\]
target/riscv: Rationalize XLEN
and operand length](#)



2021-10-07

V8 UXLEN Alistair
merged

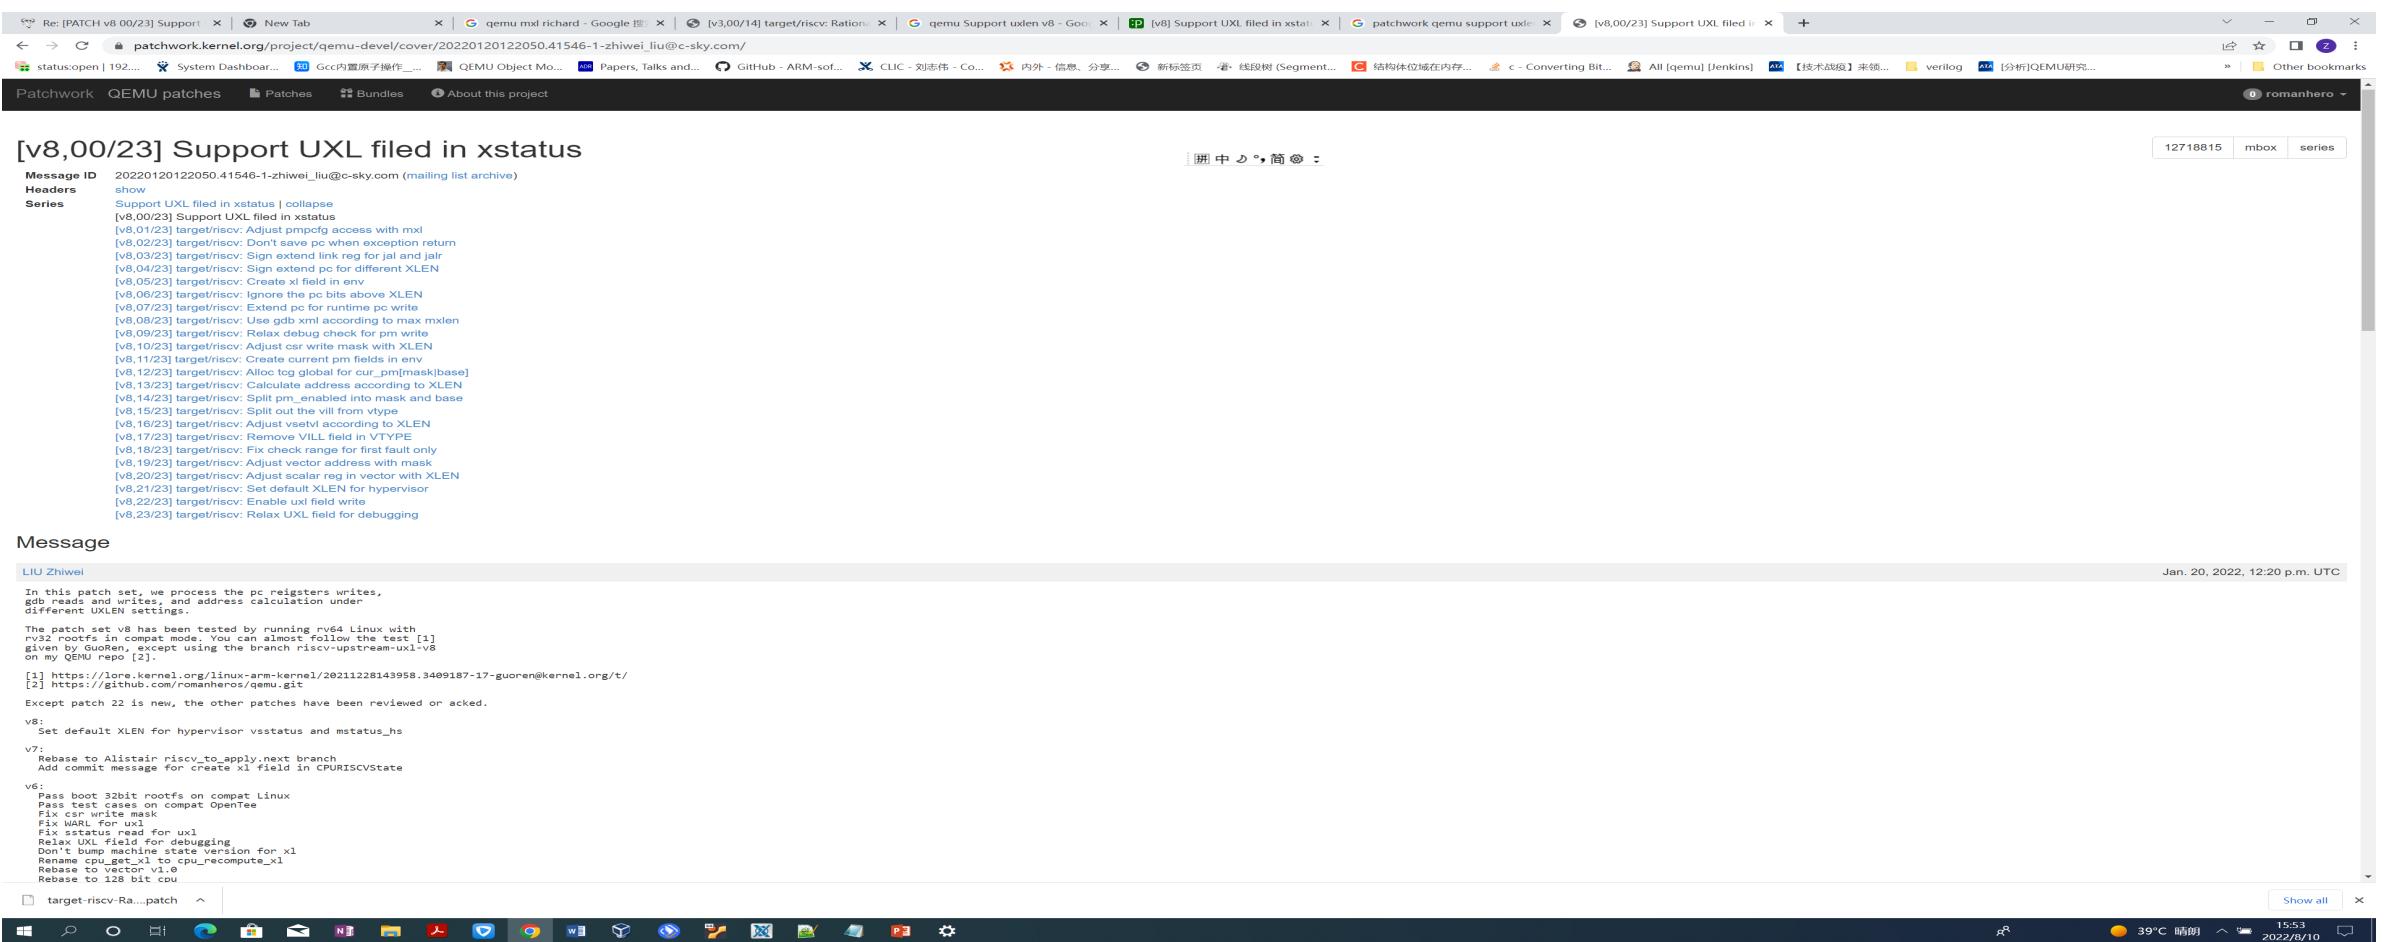
[\[PATCH v8 00/23\] Support
UXL filed in xstatus](#)



2022-1-21

Status

- The v8 UXLEN support patch set(23 patch) has been merged into QEMU v7.0.0
- Both Richard MXL and Frédéric Pétrot RV128 have been merged.
- Dynamic UXLEN support is ready.
- Neither dynamic SXLEN nor MXLEN is supported. Once both supported, we can remove RV32 system mode QEMU.



RV64 LINUX COMPAT for 32-bit U-mode

03

Linux-5.19 RISC-V statics



ARTICLES & REVIEWS NEWS ARCHIVE FORUMS PREMIUM CATEGORIES

RISC-V With Linux 5.19 Allows Running RV32 32-bit Binaries On RV64, Adds Svpbmt

Written by Michael Larabel in RISC-V on 1 June 2022 at 03:00 AM EDT. 10 Comments



On Tuesday the RISC-V architecture changes were merged into the in-development Linux 5.19 kernel with several new features in tow.

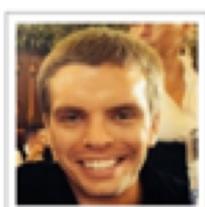
As previously covered, RISC-V with Linux 5.19 brings [the new "compat" subsystem for running 32-bit binaries on 64-bit kernels](#). This RV32 on RV64 support is interesting since the Linux kernel has always catered to 64-bit RISC-V and not much in the way of modern RV32 usage I am aware of at least where someone would be chasing mainline Linux usage.

The other big change on the RISC-V front this cycle is Svpbmt extension support for Supervisor-Mode: Page-Based Memory Types. RISC-V's Svpbmt allows for memory attributes to be encoded directly in pages. See the [riscv-isa-manual](#) for more details.

The RISC-V code for Linux 5.19 also adds support for kexec_file_load as the newer Kexec system call that is file-based and relies on FDs being added for the kernel and initramfs. There is also a new ticket-based spinlock system and other smaller fixes.

More details on the RISC-V changes premiering in Linux 5.19 this summer via [this Git merge](#).

About The Author



Michael Larabel is the principal author of Phoronix.com and founded the site in 2004 with a focus on enriching the Linux hardware experience. Michael has written more than 20,000 articles covering the state of Linux hardware support, Linux performance, graphics drivers, and other topics. Michael is also the lead developer of the Phoronix Test Suite, Phoromatic, and OpenBenchmarking.org automated benchmarking software. He can be followed via [Twitter](#), [LinkedIn](#), or contacted via [MichaelLarabel.com](#).

Linux-5.19 two important features: COMPAT and Svpbmt

https://kernelnewbies.org/Linux_5.19#RISCV
<https://www.cnx-software.com/2022/08/01/linux-5-19-release-main-changes-arm-risc-v-and-mips-architectures/>
<https://9to5linux.com/linux-kernel-5-19-officially-released-this-is-whats-new/>
<https://tinylab.org/tinylab-weekly-6-1st-2022/>

Guo Ren (21) // T-HEAD
 Heiko Stuebner (16) // VRULL & T-HEAD
 Palmer Dabbelt (14) // Rivos
 Conor Dooley (11) // Microchip
 Anup Patel (10) // ventana
 Li Zhengyu (5) // Huawei
 Christoph Hellwig (3) // ???
 Alexandre Ghti (3) // bootlin
 Liao Chang (3) // Huawei
 Nathan Chancellor (2) // ???
 Zong Li (2) // Sifive
 Xianting Tian (1)
 Chuanhua Han (1)
 Niklas Cassel (1)
 Uwe Kleine-König (1)
 Jisheng Zhang (1)
 Masahiro Yamada (1)
 Samuel Holland (1)
 Tobias Klauser (1)
 Ben Dooks (1)
 Krzysztof Kozlowski (1)
 Emil Renner Berthing (1)
 Julia Lawall (1)
 Lukas Bulwahn (1)
 Alexandre Ghti (1)
 Atish Patra (1)
 Jiapeng Chong (1)

27 contributors from various organizations contributed 106 RISC-V patches. 43 of them are directly related to the XuanTie CPU. In addition, Huawei (6) and T-HEAD (1) have improved the Kdump/Kexec feature to prepare for the crash utils upstream.

Heiko Stuebner (16):
 riscv: integrate alternatives b
 riscv: allow different stages w
 riscv: implement module alterna
 riscv: implement ALTERNATIVE_2
 riscv: extend concatenated alte
 riscv: prevent compressed instr
 riscv: move boot alternatives t
 riscv: Fix accessing pfn bits i
 riscv: add RISC-V Svpbmt extens
 riscv: remove FIXMAP_PAGE_IO an
 riscv: don't use global static
 riscv: add memory-type errata f
 riscv: fix dependency for t-hea
 riscv: drop cpufeature_apply_fe
 riscv: Improve description for
 riscv: don't warn for sifive er
 Nathan Chancellor (1):
 riscv: Fix ALT_THEAD_PMA's asm

The Svpbmt feature contains 17 patches from two contributors. It makes D1's XuanTie CPU could run upstream kernel.

Christoph Hellwig (3):
 uapi: simplify __ARCH_FLOCK{,64}_PAD a littl
 uapi: always define F_GETLK64/F_SETLK64/F_SE
 compat: consolidate the compat_flock{,64} de
 Guo Ren (21):
 arch: Add SYSVIPC_COMPAT for all architectur
 fs: stat: compat: Add __ARCH_WANT_COMPAT_STA
 asm-generic: compat: Cleanup duplicate defin
 syscalls: compat: Fix the missing part for -
 riscv: Fixup difference with defconfig
 riscv: compat: Add basic compat data type im
 riscv: compat: Support TASK_SIZE for compat
 riscv: compat: syscall: Add compat_sys_call_
 riscv: compat: syscall: Add entry.S implemen
 riscv: compat: process: Add UXL_32 support i
 riscv: compat: Add elf.h implementation
 riscv: compat: Add hw capability check for e
 riscv: compat: vdso: Add COMPAT_VDSO base co
 riscv: compat: vdso: Add setup additional pa
 riscv: compat: signal: Add rt_frame implemen
 riscv: compat: ptrace: Add compat_arch_ptrac
 riscv: compat: Add COMPAT_Kbuild skeletal su
 riscv: atomic: Cleanup unnecessary definitio
 riscv: atomic: Optimize dec_if_positive fund
 riscv: atomic: Add custom conditional atomic
 riscv: compat: Using seperated vdso_maps for
 Palmer Dabbelt (1):
 RISC-V: Add support for rv32 userspace via C
 Emil Renner Berthing (1):
 riscv: compat: vdso: Fix vdso_install target

The COMPAT feature contains 26 patches from 4 contributors.

COMPAT Introduction (V1~V9, cross two dev cycles)

Currently, in the Linux 64-bit architecture, x86, parisc, powerpc, arm64, s390, mips, and sparc all support COMPAT, and RISC-V has become the 8th architecture to support this feature. Different from its predecessors, RISC-V is the first to implement COMPAT based on asm-generic/unistd.h, and it is also the most standard. Here is work list:

- Cleanup syscall common code with community maintainers (Thx Arnd, Christoph)
- Create the base compat data structure
- Implement double TASK_SIZE
- Implement compat_sys_call_table
- Implement compat_syscall exception handler in entry.S
- Implement double ELF detect
- Implement compat_vdso
- Implement compat_signal
- Implement compat_ptrace

Agenda

- HW cost
- RV32 Linux kernel needs contributions
- eBPF JIT problem
- ELF size comparison
- Why not RV64 ILP32 ABI for Linux?

HW cost

U32 (Cheap)

- Zero extension or truncate
- Sign-extend BIT[31] to 64-bit GPR (x1~x31)
- Move between GPR and FPR/VPR

S32+M32 (Expensive)

- Sv32
- CSR xlen 32/64
- CSR h/l for 32-bit
- PMP, timer, CLIT/PLIT, PMU ...

So, COMPAT is little in HW cost for an existed 64-bit core.

RV32 Linux kernel needs contributions

ch/riscv/Kconfig:

```
ect HAVE_GENERIC_VDSO if MMU && 64BIT
ect GENERIC_TIME_VSYSCALL if MMU && 64BIT
ect GENERIC_ATOMIC64 if !64BIT
ect HAVE_ARCH_KASAN if MMU && 64BIT
ect HAVE_ARCH_KASAN_VMALLOC if MMU && 64BIT
ect HAVE_ARCH_KFENCE if MMU && 64BIT
ect HAVE_ARCH_VMAP_STACK if MMU && 64BIT
ect HAVE_ARCH_TRANSPARENT_HUGE PAGE if 64BIT && MMU
ect ARCH_WANT_HUGE_PMD_SHARE if 64BIT
ect GENERIC_BUG_RELATIVE_POINTERS if 64BIT
```

```
hfig ARCH_RV32I
    bool "RV32I"
    depends on NONPORTABLE
    select 32BIT
    select GENERIC_LIB_ASHLDI3
    select GENERIC_LIB_ASHRDI3
    select GENERIC_LIB_LSHRDI3
    select GENERIC_LIB_UCMPDI2
    select MMU
```

```
hfig ARCH_RV64I
    bool "RV64I"
    select 64BIT
    select ARCH_SUPPORTS_INT128 if CC_HAS_INT128
    select HAVE_DYNAMIC_FTRACE
    select HAVE_DYNAMIC_FTRACE_WITH_REGS
    select HAVE_FTRACE_MCOUNT_RECORD if !XIP_KERNEL
    select HAVE_FUNCTION_GRAPH_TRACER
    select HAVE_FUNCTION_TRACER if !XIP_KERNEL
    select SWIOTLB if MMU
```

```
hfig RISCV_ISA_SVPBMT // 5.19
    bool "SVPBMT extension support"
    depends on 64BIT && MMU
```

```
hfig KEXEC_FILE // 5.19
    bool "kexec file based system call"
    ...
```

JavaScript RISC-V Instruction Set Architecture ...

RISC-V Mentorship: Porting V8 to RISC-V R32G

Terms

Spring 2022

The RISC-V Mentorship Program enables one or more 12-week internship-style projects per session, funded by RISC-V, to match mentors/project leaders together with mentees/interns. Mentees are guided through a series of milestones by one or more project mentors, with whom the mentees meet on a weekly basis.

This program pairs one mentee with an experienced mentor to deliver a V8 JavaScript engine port for a 32-bit RISC-V core.

The V8 JavaScript engine for RISCV64G has been upstreamed to Chromium recently. As a basic component for the Chromium web browser and node.js, it would enlarge RISCV's application scenario. Although RV32G V8 port would be quite similar to RV64G V8 port, it is still in the TODO list. Porting and enable the RV32G on V8 will bring the embedded RISCV software ecosystem more applications, make RISC-V embed processors more competitive.

Congrats! The Google V8 JavaScript Engine now supports RV32GC.



Wei Wu (吴伟)

Hi, everyone

We are glad to share that the RV32GC porting of V8 (the JavaScript engine of the Chrome browser and Node.js) is merged by Google V8 upstream today[1].

The porting has been carried out in a cooperative way as a mentorship project of RVI and the Linux Foundation[2]. Thanks to the communication channel provided by the mentorship project (special thanks to Megan!), people who are willing to contribute to RISC-V could join together more effectively and easily.

Now V8 can support both 32bit and 64bit backends for RISC-V ISA. The team from PLCT will maintain it for the long term. Reporting bugs at GitHub repo site[3] is welcome. We hope this would be a stepping stone to enable more applications for the 32bit RISC-V world.

- [1] <https://chromium-review.googlesource.com/c/v8/v8/+/3736732>
- [2] <https://mentorship.lfx.linuxfoundation.org/project/2021e650-c533-4671-afed-bf87c089af09>
- [3] <https://github.com/riscv-collab/v8>

--
Best wishes,
Wei Wu (吴伟)
PLCT Lab, ISCAS

8月25日周四 · 下午 · 分会场B

16:30

V8 RV32GC 移植项目进展介绍

邱吉
中科院软件所

17:00

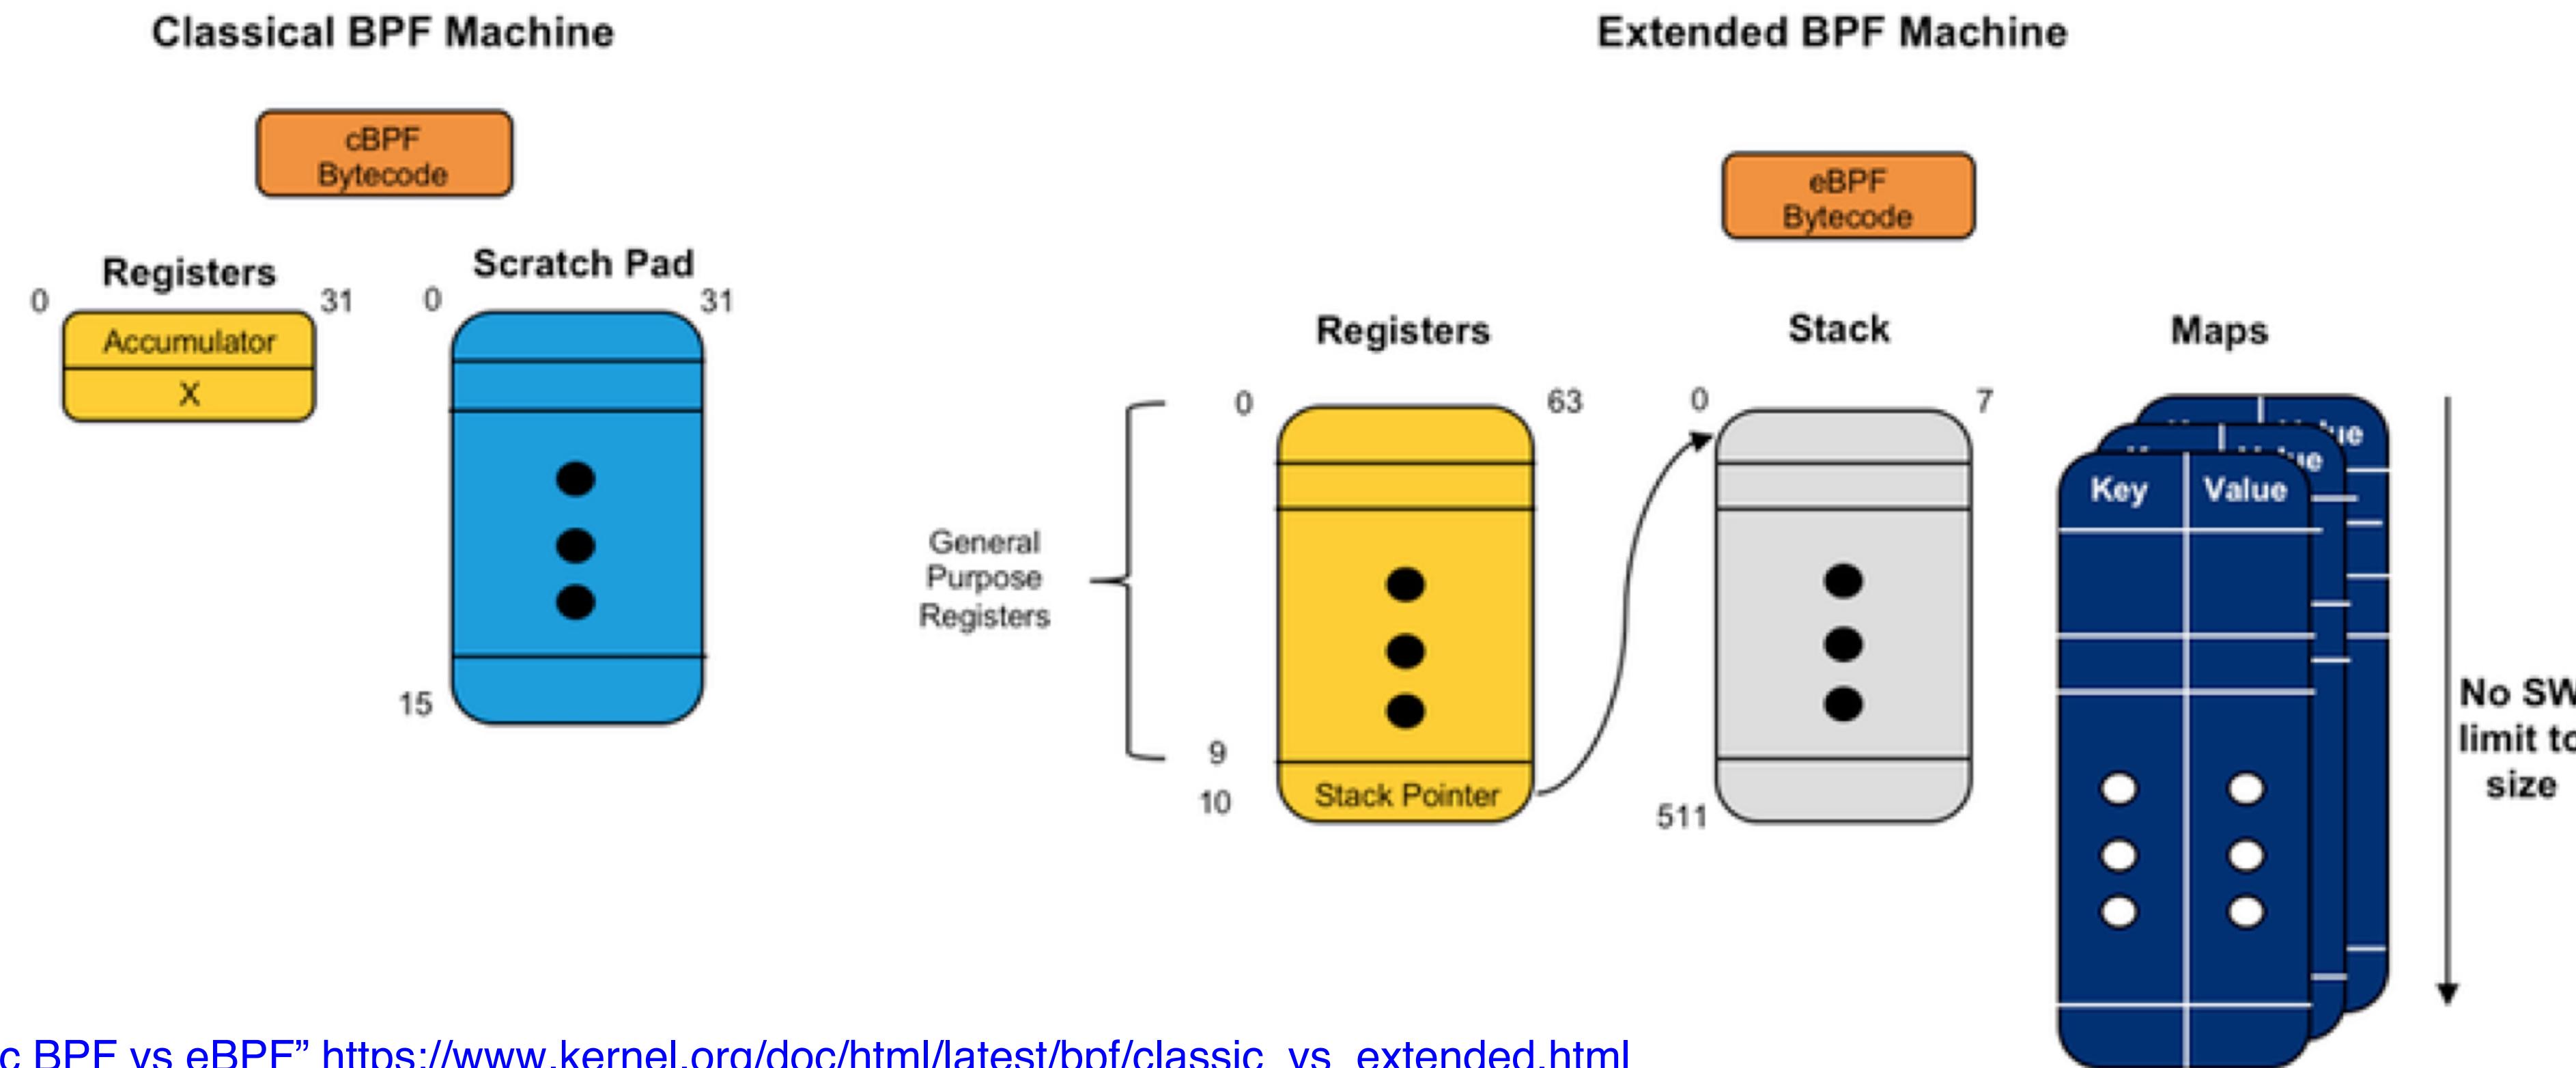
OpenJDK的RV32G支持

史宁宁
中科院软件所PLCT实验室
 RISC-V 中国峰会 2022

eBPF JIT problem

- Number of registers increase from 2 to 10
- Register width increases from 32-bit to 64-bit
- Conditional jt/jf targets replaced with jt/fall-through
- Introduces bpf_call insn and register passing convention for zero overhead calls from/to other kernel functions

rv32-Linux with cBPF
rv64-Linux with eBPF



"Classic BPF vs eBPF" https://www.kernel.org/doc/html/latest/bpf/classic_vs_extended.html

ELF size comparison

64kernel + 32rootfs

```
[ 0.000000] Linux version 5.16.0-rc6-00017-g750f87086bdd-dirty (guoren@guoren-z87-HD3) (riscv64-unknown-linux-gnu-gcc (GCC) 10.2.0, GNU ld (GNU Binutils) 2.37) #96 SMP Tue Dec 28 21:01:55 CST 2021
[ 0.000000] OF: fdt: Ignoring memory range 0x80000000 - 0x80200000
[ 0.000000] Machine model: riscv-virtio,gemu
[ 0.000000] earlycon: sbi0 at I/O port 0x0 (options '')
[ 0.000000] printk: bootconsole [sbi0] enabled
[ 0.000000] efi: UEFI not found.
[ 0.000000] Zone ranges:
[ 0.000000]   DMA32   [mem 0x0000000080200000-0x0000000083fffff]
[ 0.000000]   Normal   empty
[ 0.000000] Movable zone start for each node
[ 0.000000] Early memory node ranges
[ 0.000000]   node 0: [mem 0x0000000080200000-0x0000000083fffff]
[ 0.000000] Initmem setup node 0 [mem 0x0000000080200000-0x0000000083fffff]
[ 0.000000] SBI specification v0.2 detected
[ 0.000000] SBI implementation ID=0x1 Version=0x9
[ 0.000000] SBI TIME extension detected
[ 0.000000] SBI IPI extension detected
[ 0.000000] SBI RFENCE extension detected
[ 0.000000] SBI v0.2 HSM extension detected
[ 0.000000] riscv: ISA extensions acdfhimsu
[ 0.000000] riscv: ELF capabilities acdfim
[ 0.000000] percpu: Embedded 17 pages/cpu s30696 r8192 d30744 u69632
[ 0.000000] Built 1 zonelists, mobility grouping on. Total pages: 15655
[ 0.000000] Kernel command line: rootwait root=/dev/vda ro console=ttyS0 earlycon=sbi
[ 0.000000] Dentry cache hash table entries: 8192 (order: 4, 65536 bytes, linear)
[ 0.000000] Inode-cache hash table entries: 4096 (order: 3, 32768 bytes, linear)
[ 0.000000] mem auto-init: stack:off, heap alloc:off, heap free:off
[ 0.000000] Virtual kernel memory layout:
[ 0.000000]   fixmap : 0xfffffffcef000000 - 0xfffffffcef000000 (2048 kB)
[ 0.000000]   pci io : 0xfffffffcef000000 - 0xfffffffcef00000000 ( 16 MB)
[ 0.000000]   vmemmap : 0xfffffffcef000000 - 0xfffffffcef00000000 (4095 MB)
[ 0.000000]   vmalloc : 0xffffffffd0000000 - 0xffffffffdfffff0000 (65535 MB)
[ 0.000000]   lowmem : 0xfffffffffe00000000 - 0xfffffffffe003e00000 ( 62 MB)
[ 0.000000]   kernel : 0xfffffffff8000000 - 0xfffffffff800000000 (2047 MB)
[ 0.000000] Memory: 52788K/63488K available (6184K kernel code, 888K rwdata, 1917K
rodata, 294K init, 297K bss, 10700K reserved, 0K cma-reserved)
[ 0.000000] SLUB: HWalign=64, Order=0-3, MinObjects=0, CPUs=1, Nodes=1
[ 0.000000] rcu: Hierarchical RCU implementation.
[ 0.000000] rcu:   RCU restricting CPUs from NR_CPUS=8 to nr_cpu_ids=1.
[ 0.000000] rcu:   RCU debug extended QS entry/exit.
[ 0.000000] Tracing variant of Tasks RCU enabled.
[ 0.000000] rcu: RCU calculated value of scheduler-enlistment delay is 25 jiffies.
[ 0.000000] rcu: Adjusting geometry for rcu_fanout_leaf=16, nr_cpu_ids=1
[ 0.000000] NR_IRQS: 64, nr_irqs: 64, preallocated irqs: 0
[ 0.000000] riscv-intc: 64 local interrupts mapped
[ 0.000000] plic: plic@c00000: mapped 53 interrupts with 1 handlers for 2 contexts.
...
Welcome to Buildroot
```

64: Memory: 52788K/63488K available (6184K kernel code, 888K rwdata, 1917K rodata, 294K init, 297K bss, 10700K reserved)

32: Memory: 51924K/61440K available (6117K kernel code, 695K rwdata, 1594K rodata, 255K init, 241K bss, 9516K reserved)

result:

	1%	22%	17%	13%	19%	11%
--	----	-----	-----	-----	-----	-----

rv32 kernel runtime KernelStack, Slab... are smaller

rv64: MemTotal: 53076 kB, MemFree: 40264 kB

rv32: MemTotal: 52176 + 2048 kB, MemFree: 41012 + 2048 kB

result :

	6%
--	----

32kernel + 32rootfs

```
[ 0.000000] Linux version 5.16.0-rc6-00017-g750f87086bdd-dirty (guoren@guoren-z87-HD3) (riscv32-buildroot-linux-gnu-gcc.br_real (Buildroot 2021.11-201-g7600ca7960-dirty) 10.3.0, GNU ld (GNU Binutils) 2.36.1) #7 SMP Tue Dec 28 21:02:21 CST 2021
[ 0.000000] OF: fdt: Ignoring memory range 0x80000000 - 0x80400000
[ 0.000000] Machine model: riscv-virtio,gemu
[ 0.000000] earlycon: sbi0 at I/O port 0x0 (options '')
[ 0.000000] printk: bootconsole [sbi0] enabled
[ 0.000000] efi: UEFI not found.
[ 0.000000] Zone ranges:
[ 0.000000]   Normal   [mem 0x0000000080400000-0x0000000083fffff]
[ 0.000000]   Movable zone start for each node
[ 0.000000] Early memory node ranges
[ 0.000000]   node 0: [mem 0x0000000080400000-0x0000000083fffff]
[ 0.000000] Initmem setup node 0 [mem 0x0000000080400000-0x0000000083fffff]
[ 0.000000] SBI specification v0.2 detected
[ 0.000000] SBI implementation ID=0x1 Version=0x9
[ 0.000000] SBI TIME extension detected
[ 0.000000] SBI IPI extension detected
[ 0.000000] SBI RFENCE extension detected
[ 0.000000] SBI v0.2 HSM extension detected
[ 0.000000] riscv: ISA extensions acdfhimsu
[ 0.000000] riscv: ELF capabilities acdfim
[ 0.000000] percpu: Embedded 12 pages/cpu s16600 r8192 d24360 u49152
[ 0.000000] Built 1 zonelists, mobility grouping on. Total pages: 15240
[ 0.000000] Kernel command line: rootwait root=/dev/vda ro console=ttyS0 earlycon=sbi
[ 0.000000] Dentry cache hash table entries: 8192 (order: 3, 32768 bytes, linear)
[ 0.000000] Inode-cache hash table entries: 4096 (order: 2, 16384 bytes, linear)
[ 0.000000] mem auto-init: stack:off, heap alloc:off, heap free:off
[ 0.000000] Virtual kernel memory layout:
[ 0.000000]   fixmap : 0x9dc00000 - 0x9e000000 (4096 kB)
[ 0.000000]   pci io : 0x9e000000 - 0x9f000000 ( 16 MB)
[ 0.000000]   vmemmap : 0x9f000000 - 0x9fffff0000 ( 15 MB)
[ 0.000000]   vmalloc : 0xa0000000 - 0xbfffffff ( 511 MB)
[ 0.000000]   lowmem : 0xc0000000 - 0xc3c00000 ( 60 MB)
[ 0.000000]   kernel : 0xc3c00000 - 0xd0000000 (2047 MB)
[ 0.000000] Memory: 51924K/61440K available (6117K kernel code, 695K rwdata, 1594K
rodata, 255K init, 241K bss, 9516K reserved, OK cma-reserved)
[ 0.000000] SLUB: HWalign=64, Order=0-3, MinObjects=0, CPUs=1, Nodes=1
[ 0.000000] rcu: Hierarchical RCU implementation.
[ 0.000000] rcu:   RCU restricting CPUs from NR_CPUS=8 to nr_cpu_ids=1.
[ 0.000000] rcu:   RCU debug extended QS entry/exit.
[ 0.000000] Tracing variant of Tasks RCU enabled.
[ 0.000000] rcu: RCU calculated value of scheduler-enlistment delay is 25 jiffies.
[ 0.000000] rcu: Adjusting geometry for rcu_fanout_leaf=16, nr_cpu_ids=1
[ 0.000000] NR_IRQS: 64, nr_irqs: 64, preallocated irqs: 0
[ 0.000000] riscv-intc: 32 local interrupts mapped
[ 0.000000] plic: plic@c00000: mapped 53 interrupts with 1 handlers for 2 contexts.
...
Welcome to Buildroot
```

Why not RV64 ILP32 ABI for Linux?

Question answered by software guys.

```

Subject: Re: [PATCH 00/13] riscv: compat: Add COMPAT mode support for rv64
Date: Mon, 27 Dec 2021 02:29:02 +0000 [thread overview]
Message-ID: <AA7091EA-C3AF-47CE-B0C5-E2ECE4133D09@jrtc27.com> (raw)
In-Reply-To: <CAJF2gTQ04uty0c9=q9Tu92togGuuygKqg3tWwfPBcuyTfLh2SQ@mail.gmail.com>

On 27 Dec 2021, at 01:16, Guo Ren <guoren@kernel.org> wrote:
>
> On Mon, Dec 27, 2021 at 4:31 AM Arnd Bergmann <arnd@arndb.de> wrote:
>>
>> On Sun, Dec 26, 2021 at 7:38 AM Guo Ren <guoren@kernel.org> wrote:
>>> On Sun, Dec 26, 2021 at 4:36 PM Jisheng Zhang <jszhang3@mail.ustc.edu.cn> wrote:
>>>> On Wed, 22 Dec 2021 20:59:30 +0800 Guo Ren <guoren@kernel.org> wrote:
>>>>> On Wed, Dec 22, 2021 at 2:10 AM Arnd Bergmann <arnd@arndb.de> wrote:
>>>>>
>>>> What about adding RV64 ILP32 support instead? This don't need HW side
>>>> modifications so can benefit all RV64.
>>>
>>> ILP32 is another topic in C Language Data Type Models and it couldn't
>>> replace the standard rv32 ecosystem.
>>> COMPAT is a common framework in Linux (7 arches have been supported),
>>> so let rv64 support COMPAT mode is considerable.
>>>
>>> Customers would choose ILP32 / RV32-compat by themself and that
>>> depends on which one has a better ecosystem.
>>>
>>> From a kernel perspective, supporting both is not much more work than
>>> supporting either of them. We had the same debate for Arm64, and ended
>>> up never merging the ILP32 patches despite them being well written
>>> and maintainable, to limit the number of supported user space ABIs
>>> as well as the possible attack vectors when there is an exploitable
>>> bug that is specific to an ABI.
>>>
>>> arm64 does support big-endian mode, which is a similar niche, but it
>>> can't easily be removed after it's already supported. Supporting normal
>>> compat mode is the easiest here because it doesn't add another user
>>> space ABI, but I'd strongly recommend not to add any other ones.
>>>
>> @Palmer Dabbelt How do you think about supporting ILP32 & COMPAT both
> in rv64? And let users vote by foot which is better.

As psABI TG co-chair I really do not want an ILP32 RV64 to exist if it
can at all be avoided. Every single attempt at an ILP32 ABI for a
64-bit architecture has failed to take off in the past, so I struggle
to see why RV64 will be any different. So, in my opinion, there is a
relatively high barrier to entry for it to be an official frozen ABI,
and without it being that I doubt upstreams will want to go near it, be
it Linux, GCC, binutils or GCC, but they can speak for themselves if
they feel otherwise.

Also, with every year that goes by, ILP32 becomes more and more limited
in what you can use it for, due to increased memory footprints...

```

Q: Why not RV64 ILP32? How about apple watch's aarch64-ILP32?

ILP32 for ARM64

From: Yury Norov <ynorov@kernel.org>
To: <linux-arm-kernel@lists.infradead.org>, <linux-kernel@vger.kernel.org>, <linux-doc@vger.kernel.org>, Arnd Bergmann <arnd@arndb.de>, Catalin Marinas <catalin.marinas@arm.com>
Subject: [PATCH v7 00/20] ILP32 for ARM64
Date: Mon, 9 Jan 2017 16:59:37 +0530
Message-ID: <1483961397-8599-1-git-send-email-ynorov@caviumnetworks.com>
Cc: Yury Norov <ynorov@kernel.org>, Andrew Pinski <pinskia@gmail.com>, Adam Borowski <kilobyte@angband.pl>, Chris Metcalf <cmetcalf@linaro.org>, Maxim Kuvyrkov <maxim.kuvyrkov@linaro.org>, Ramana Radhakrishna Weimer <fweimer@redhat.com>, Bamvor Zhangjian <bamvor.zhangjian@huawei.com>, And Metcalf <cmetcalf@mellanox.com>, Heiko Carstens <heiko.carstens@de.ibm.com>, <schwider@kernel.org>, Joseph Myers <joseph@codesourcery.com>, <christoph.muellner@theobroma-solutions.com>, <klimov.linux@at-gmail.com>, <Nathan.Lynch@mentor.com>, <agraf@t-suse.de>, <Prasun.Ka@huawei.com>, <philipp.tomsich@theobroma-systems.com>, <manuel.montezelo@kernel.org>, <davem@davemloft.net>, <zhouchengming1@huawei.com>

This series enables aarch64 with ilp32 mode.

As supporting work, it introduces ARCH_32BIT_OFF_T configuration option that is enabled for existing 32-bit architectures but disabled for new arches (so 64-bit off_t is used by new userspace). Also it deprecates getrlimit and setrlimit syscalls prior to prlimit64.

This version is based on linux-next from 2017-01-09. It works with glibc-2.25, and tested with LTP, glibc testsuite, trinity, lmbench, CPUSSpec.

This is not RFC anymore. I believe that all ABI and implementation issues are resolved now. The way that kernel clears registers top halves is probably the last question, and because there's no objection for current approach for more than 6 month, I think, community agrees with it.

Patches 1, 2, 3 and 8 are general, and may be applied separately.

Current version does not introduce ABI changes comparing to RFC3. Kernel and GLIBC trees:
<https://github.com/ynorov/linux/tree/ilp32-2017-01-09>
<https://github.com/ynorov/glibc/tree/dev>

(GLIBC patches are managed by Steve Ellsey, so my tree is only for reference.)

Changes:
 v3: <https://lkml.org/lkml/2014/9/3/704>
 v4: <https://lkml.org/lkml/2015/4/13/691>
 v5: <https://lkml.org/lkml/2015/9/29/911>
 v6: <https://lkml.org/lkml/2016/5/23/661>
 v7: RFC nowrap: <https://lkml.org/lkml/2016/6/17/990>
 v7: RFC2 nowrap: <https://lkml.org/lkml/2016/8/17/245>
 v7: RFC3 nowrap: <https://lkml.org/lkml/2016/10/21/883>
 v7: - 32-bit off_t depreciation is splitted for compat and native 32-bit arches, as it was initially done (patches 1, 2);
 - getrlimit() and setrlimit() syscalls deprecated for aarch64/ilp32 and all new architectures;
 - documentation is cleaned up (patch 4);
 - compat-related definitions moved from aarch64/include/elf.h to binfmt_elf32.c (patch 11);
 - for ptrace, execution mode detection is performed at runtime, as it was in v4 (patch 18)

Andrew Pinski (6):
 arm64: rename COMPAT to AARCH32_ELO in Kconfig
 arm64: ensure the kernel is compiled for LP64
 arm64: uapi: set __BITS_PER_LONG correctly for ILP32 and LP64
 arm64: ilp32: add sys_ilp32.c and a separate table (in entry.S) to use it
 arm64: ilp32: introduce ilp32-specific handlers for sigframe and ucontext
 arm64: ilp32: add ARM64_ILP32 to Kconfig
 Philipp Tomsich (1):
 arm64: ilp32: add vdso-ilp32 and use for signal return
 Yury Norov (13):
 compat ABI: use non-compat openat and open_by_handle_at variants
 32-bit ABI: introduce ARCH_32BIT_OFF_T config option
 asm-generic: Drop getrlimit and setrlimit syscalls from default list
 arm64: ilp32: add documentation on the ILP32 ABI for ARM64
 thread: move thread bits accessors to separated file
 arm64: introduce is_a32_task and is_a32_thread (for AArch32 compat)
 arm64: ilp32: add is_ilp32_compat_{task,thread} and TIF_32BIT_AARCH64
 arm64: introduce binfmt_elf32.c
 arm64: ilp32: introduce binfmt_ilp32.c
 arm64: ilp32: share aarch32 syscall handlers
 arm64: signal: share lp64 signal routines to ilp32
 arm64: signal32: move ilp32 and aarch32 common code to separated file
 arm64: ptrace: handle ptrace_request differently for aarch32 and ilp32
 Documentation/arm64/ilp32.txt | 45 +++++++
 arch/Kconfig | 4 +
 arch/arc/Kconfig | 1 +
 arch/arm/Kconfig | 1 +
 arch/arm64/Kconfig | 19 ++
 arch/arm64/Makefile | 5 +
 arch/arm64/include/asm/compat.h | 19 +-
 arch/arm64/include/asm/elf.h | 32 +---
 arch/arm64/include/asm/fpsimd.h | 2 +
 arch/arm64/include/asm/trace.h | 2 +
 arch/arm64/include/asm/hwcap.h | 6 +
 arch/arm64/include/asm/is_compatible.h | 90 ++++++
 arch/arm64/include/asm/memory.h | 5 +
 arch/arm64/include/asm/processor.h | 11 +
 arch/arm64/include/asm/seccomp.h | 2 +
 arch/arm64/include/asm/signals32.h | 2 +
 arch/arm64/include/asm/signals32_common.h | 9 +
 arch/arm64/include/asm/signals_common.h | 27 +++
 arch/arm64/include/asm/signals_ilp32.h | 33 +++
 arch/arm64/include/asm/syscall.h | 38 +++++
 arch/arm64/include/asm/thread_info.h | 2 +
 arch/arm64/include/asm/unistd.h | 4 +
 arch/arm64/include/asm/vdso.h | 8 +
 arch/arm64/include/uapi/asm/bitsperlong.h | 6 +
 arch/arm64/include/uapi/asm/unistd.h | 13 +
 arch/arm64/kernel/Makefile | 18 ++
 arch/arm64/kernel/asm-offsets.c | 9 +
 arch/arm64/kernel/binfmt_elf32.c | 32 +---
 arch/arm64/kernel/binfmt_ilp32.c | 98 ++++++
 arch/arm64/kernel/cpufeature.c | 8 +
 arch/arm64/kernel/cpuinfo.c | 20 +-
 arch/arm64/kernel/entry.S | 34 +---
 arch/arm64/kernel/entry32.S | 80 -----
 arch/arm64/kernel/entry32_common.S | 107 ++++++
 arch/arm64/kernel/entry_ilp32.S | 22 +---
 arch/arm64/kernel/head.S | 2 +
 arch/arm64/kernel/hw_breakpoint.c | 8 +
 arch/arm64/kernel/perf_regs.c | 2 +
 arch/arm64/kernel/process.c | 7 +
 arch/arm64/kernel/ptrace.c | 80 ++++++
 arch/arm64/kernel/signals.c | 102 ++++++
 arch/arm64/kernel/signals32.c | 107 -----
 arch/arm64/kernel/signals32_common.c | 135 ++++++
 arch/arm64/kernel/signals_ilp32.c | 170 ++++++
 arch/arm64/kernel/sys_ilp32.c | 100 ++++++
 arch/arm64/kernel/trace.h | 2 +
 arch/arm64/kernel/vdso_ilp32/Makefile | 74 +++++
 arch/arm64/kernel/vdso_ilp32/vdso-ilp32.S | 33 +---
 arch/arm64/kernel/vdso-ilp32/vdso-ilp32.lds.S | 95 ++++++
 arch/arm64/kernel/vdso.c | 70 ++++++

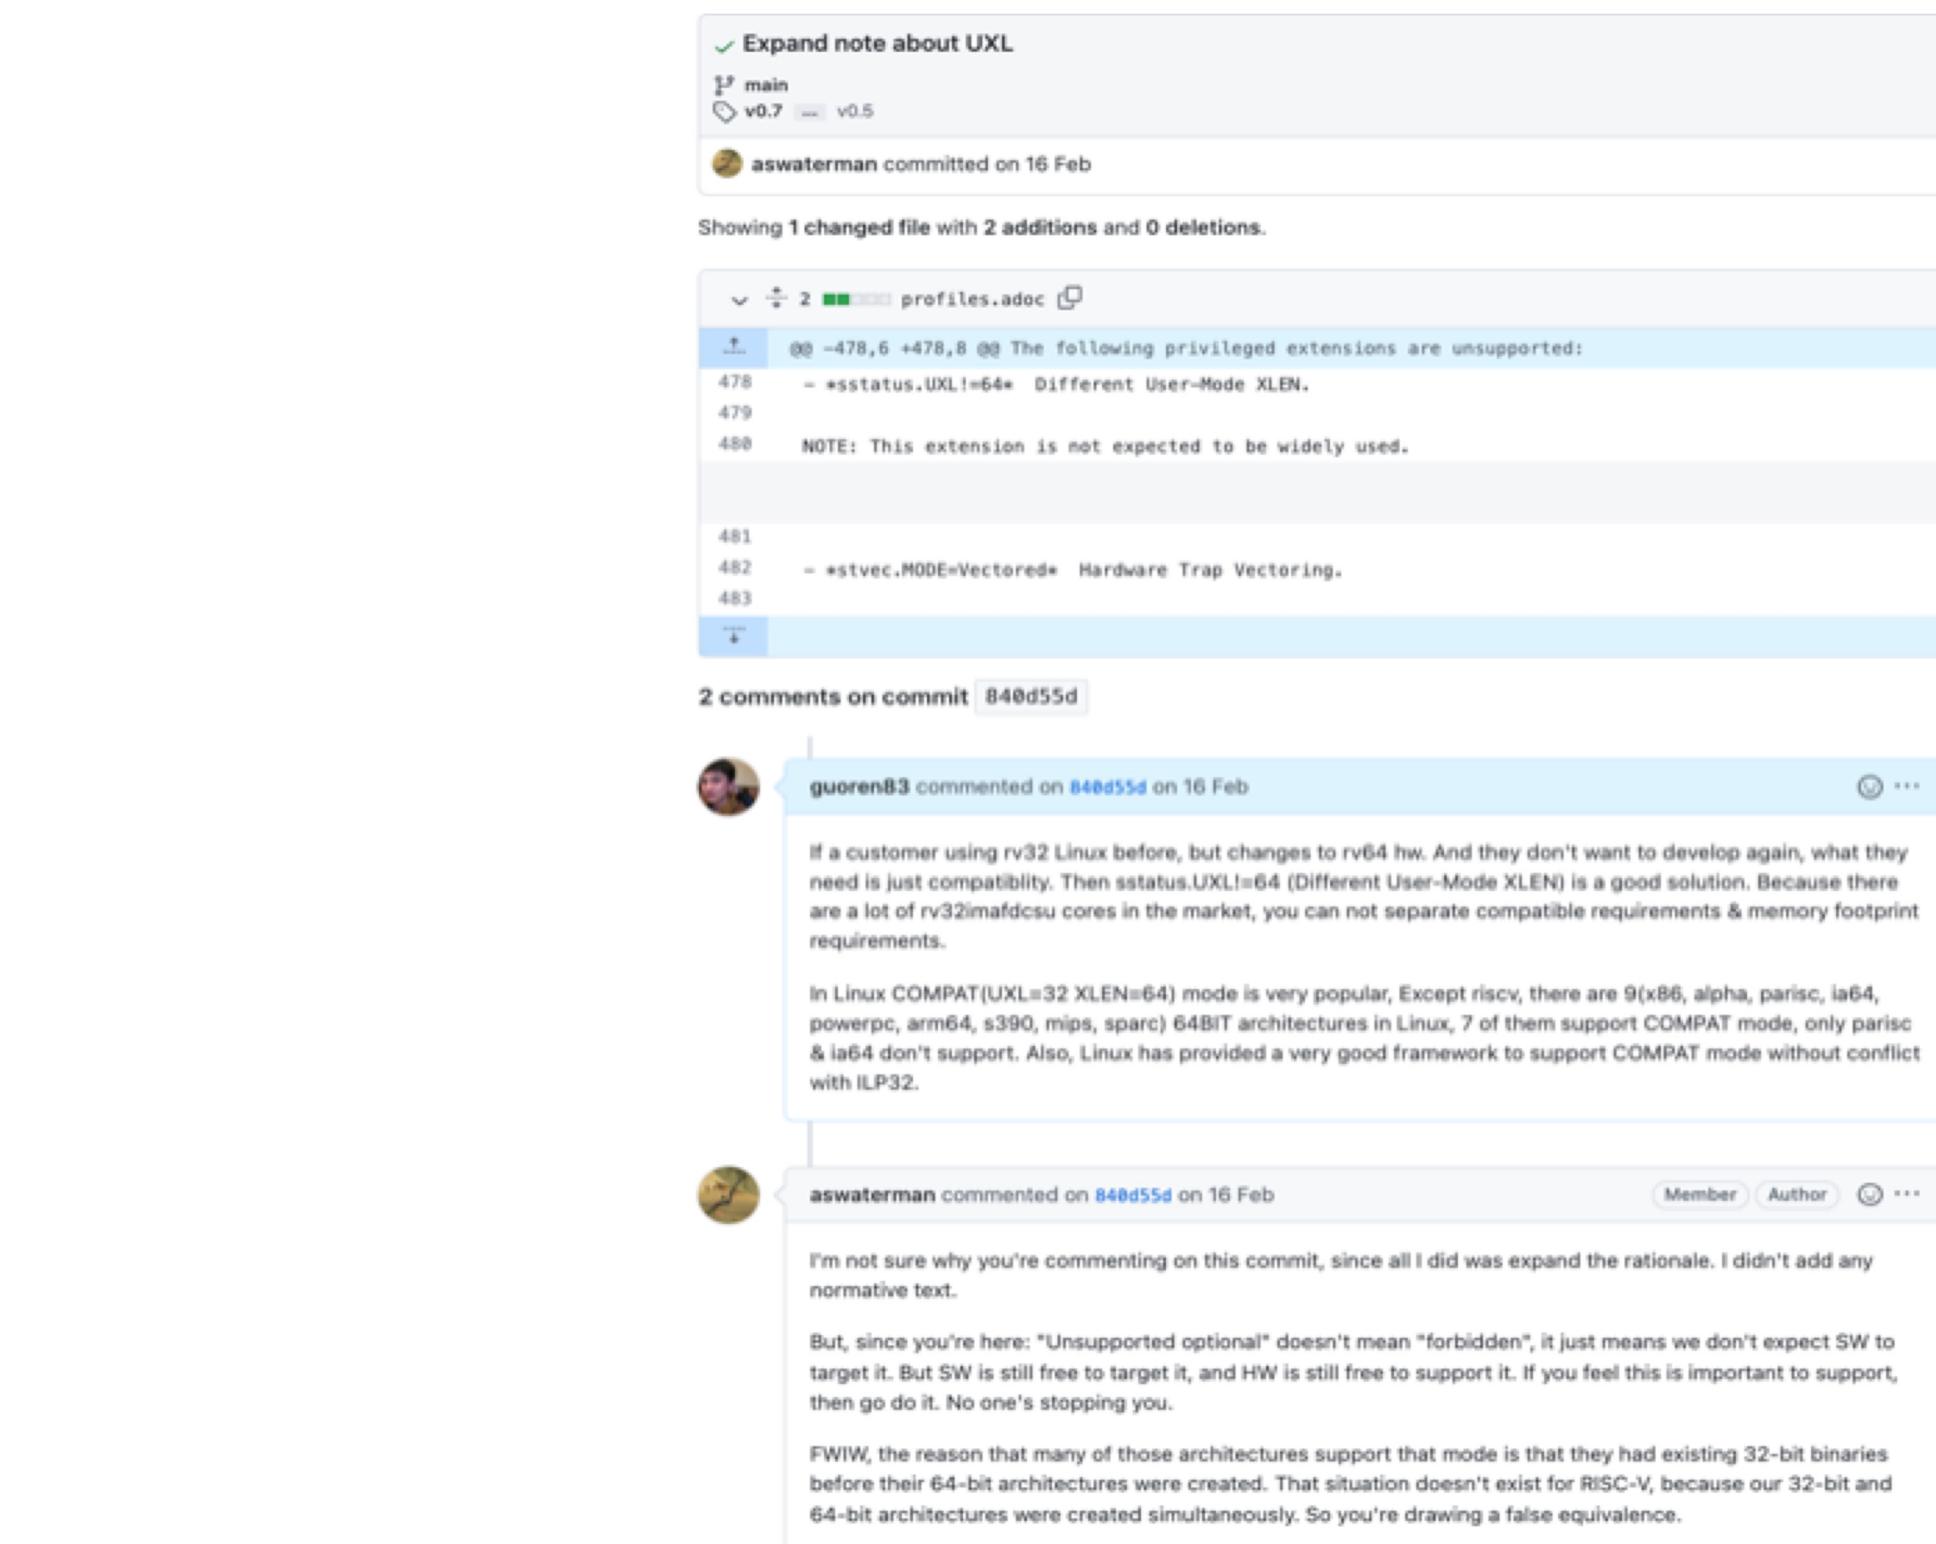
Profiles' problem discussion

5.2.4. RVA20S64 Unsupported Optional Extensions

The following privileged extensions are unsupported:

- `Ssu32xl sstatus.UXL =32`

Note This extension is not expected to be widely used, as RISC-V applications processors were initially, and are still predominantly, XLEN=64, avoiding the need to support legacy XLEN=32 binaries. Applications that want to use 32-bit pointers to reduce memory footprint should consider using a forthcoming RV64 ILP32 ABI instead.



This screenshot shows a GitHub commit page for a file named `profiles.adoc`. The commit was made by `aswaterman` on Feb 16. The commit message includes a note about unsupported extensions:

```

@@ -478,6 +478,8 @@ The following privileged extensions are unsupported:
478   *sstatus.UXL!=64* Different User-Mode XLEN.
479
480   NOTE: This extension is not expected to be widely used.

481
482   **stvec.MODE=Vectored* Hardware Trap Vectoring.
483

```

Below the commit message, there are two comments:

- guoren83** commented on Feb 16: If a customer using rv32 Linux before, but changes to rv64 hw. And they don't want to develop again, what they need is just compatibility. Then `sstatus.UXL!=64` (Different User-Mode XLEN) is a good solution. Because there are a lot of rv32imafdcu cores in the market, you can not separate compatible requirements & memory footprint requirements.
- aswaterman** commented on Feb 16: I'm not sure why you're commenting on this commit, since all I did was expand the rationale. I didn't add any normative text.

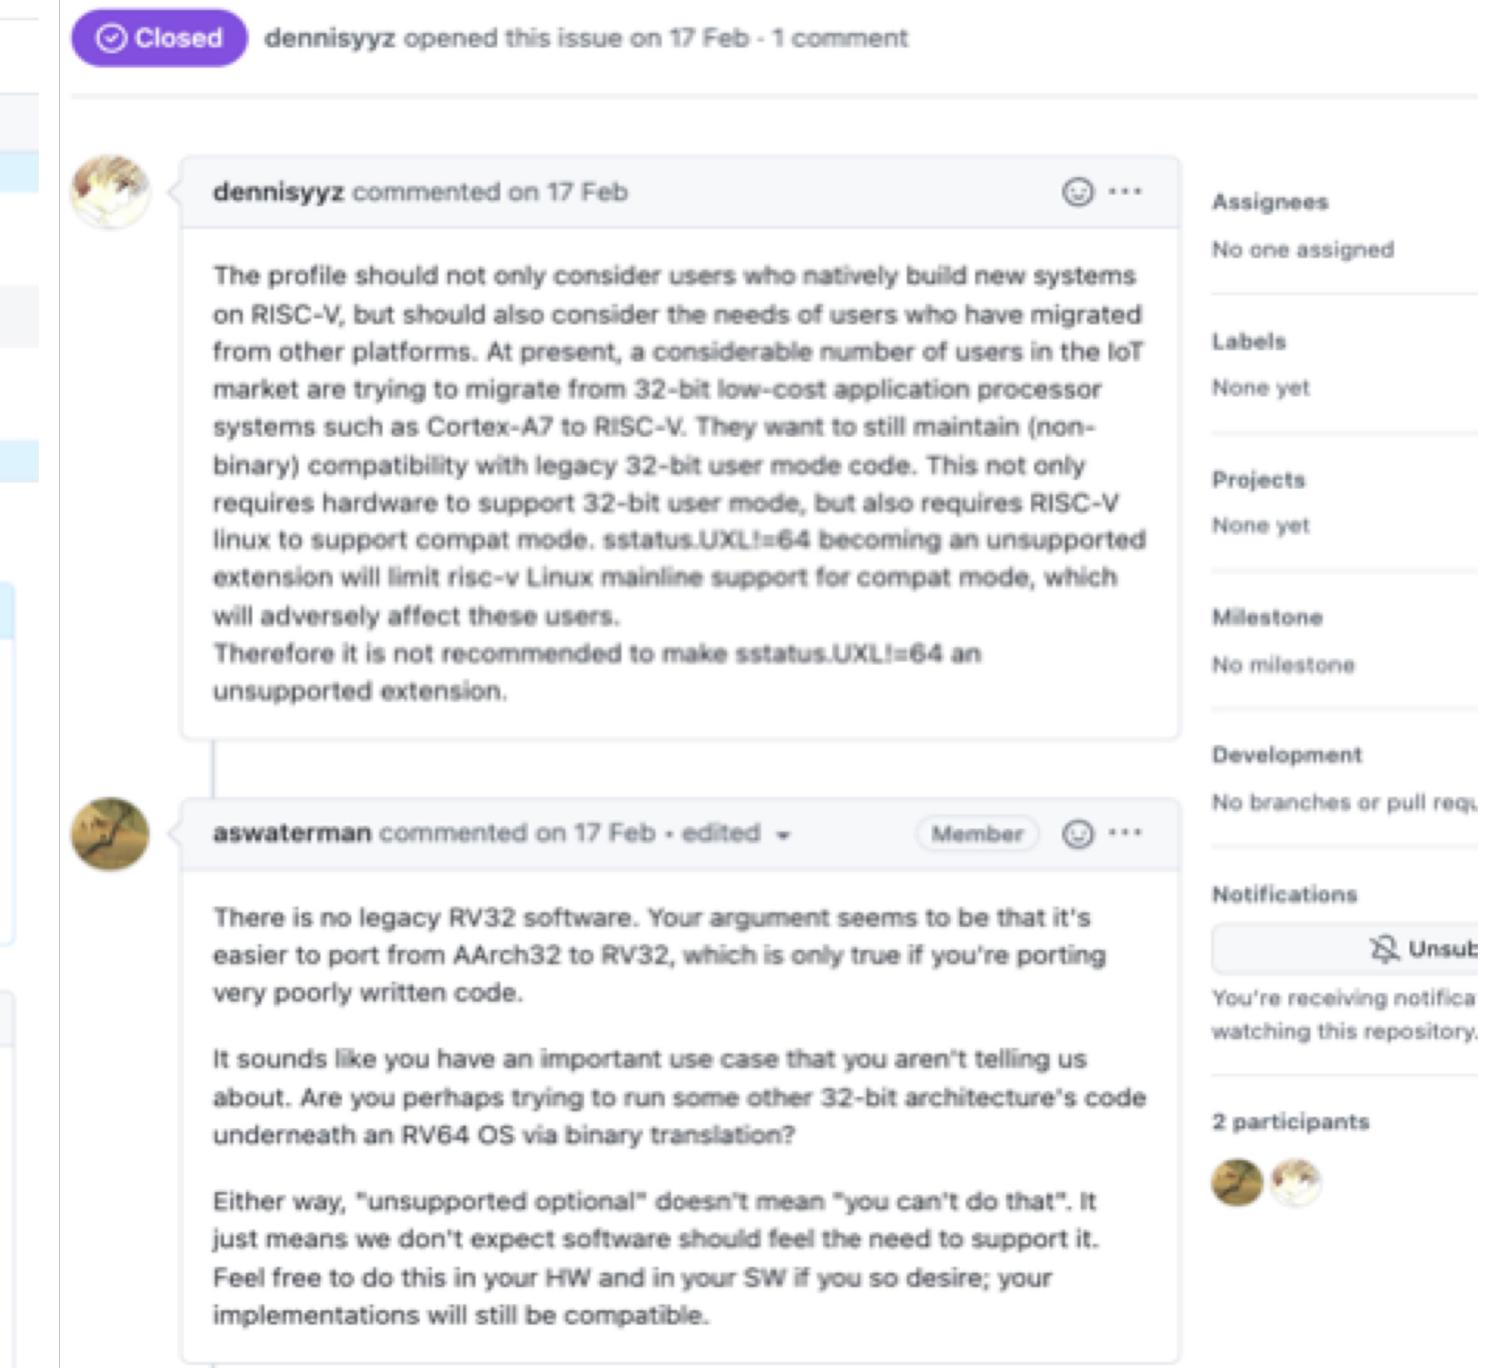
6.2.4. RVA22S64 Unsupported Optional Extensions

- `Ssu32xl sstatus.UXL =32`

Note This extension is not expected to be widely used.

<https://github.com/riscv/riscv-profiles/blob/main/profiles.adoc>

It is not recommended to make `sstatus.UXL!=64` an unsupported extension #15



This screenshot shows a GitHub issue titled "It is not recommended to make `sstatus.UXL!=64` an unsupported extension". The issue was opened by `dennisyzz` on Feb 17 and closed by `dennisyzz` on Feb 17. The issue contains the following content:

dennisyzz commented on 17 Feb

The profile should not only consider users who natively build new systems on RISC-V, but should also consider the needs of users who have migrated from other platforms. At present, a considerable number of users in the IoT market are trying to migrate from 32-bit low-cost application processor systems such as Cortex-A7 to RISC-V. They want to still maintain (non-binary) compatibility with legacy 32-bit user mode code. This not only requires hardware to support 32-bit user mode, but also requires RISC-V linux to support compat mode. `sstatus.UXL!=64` becoming an unsupported extension will limit risc-v Linux mainline support for compat mode, which will adversely affect these users. Therefore it is not recommended to make `sstatus.UXL!=64` an unsupported extension.

aswaterman commented on 17 Feb - edited

There is no legacy RV32 software. Your argument seems to be that it's easier to port from AArch32 to RV32, which is only true if you're porting very poorly written code.

It sounds like you have an important use case that you aren't telling us about. Are you perhaps trying to run some other 32-bit architecture's code underneath an RV64 OS via binary translation?

Either way, "unsupported optional" doesn't mean "you can't do that". It just means we don't expect software to feel the need to support it. Feel free to do this in your HW and in your SW if you so desire; your implementations will still be compatible.

Profiles' principle

Profiles are ...

- A mechanism to identify a group of extensions (in behavior, state) that work together
- Like macros. Profiles consist of Extensions. Extensions consist of a group of rules, state, and behaviors of extensions
- What we ask the upstream projects (gcc, binutils, and Linux distros to target. Major and minor releases
- • The way we reduce fragmentation in the RISC-V community



Some suggestions at the end

- Focus on native ILP32 in Linux userspace
- Choose RV64 processors with COMPAT
- contribute rv32-Linux actively

谢谢

THANK YOU

阿里云



平头哥