

IOMMU for Xuantie-based Platforms

Siqi Zhao
Aug, 2022

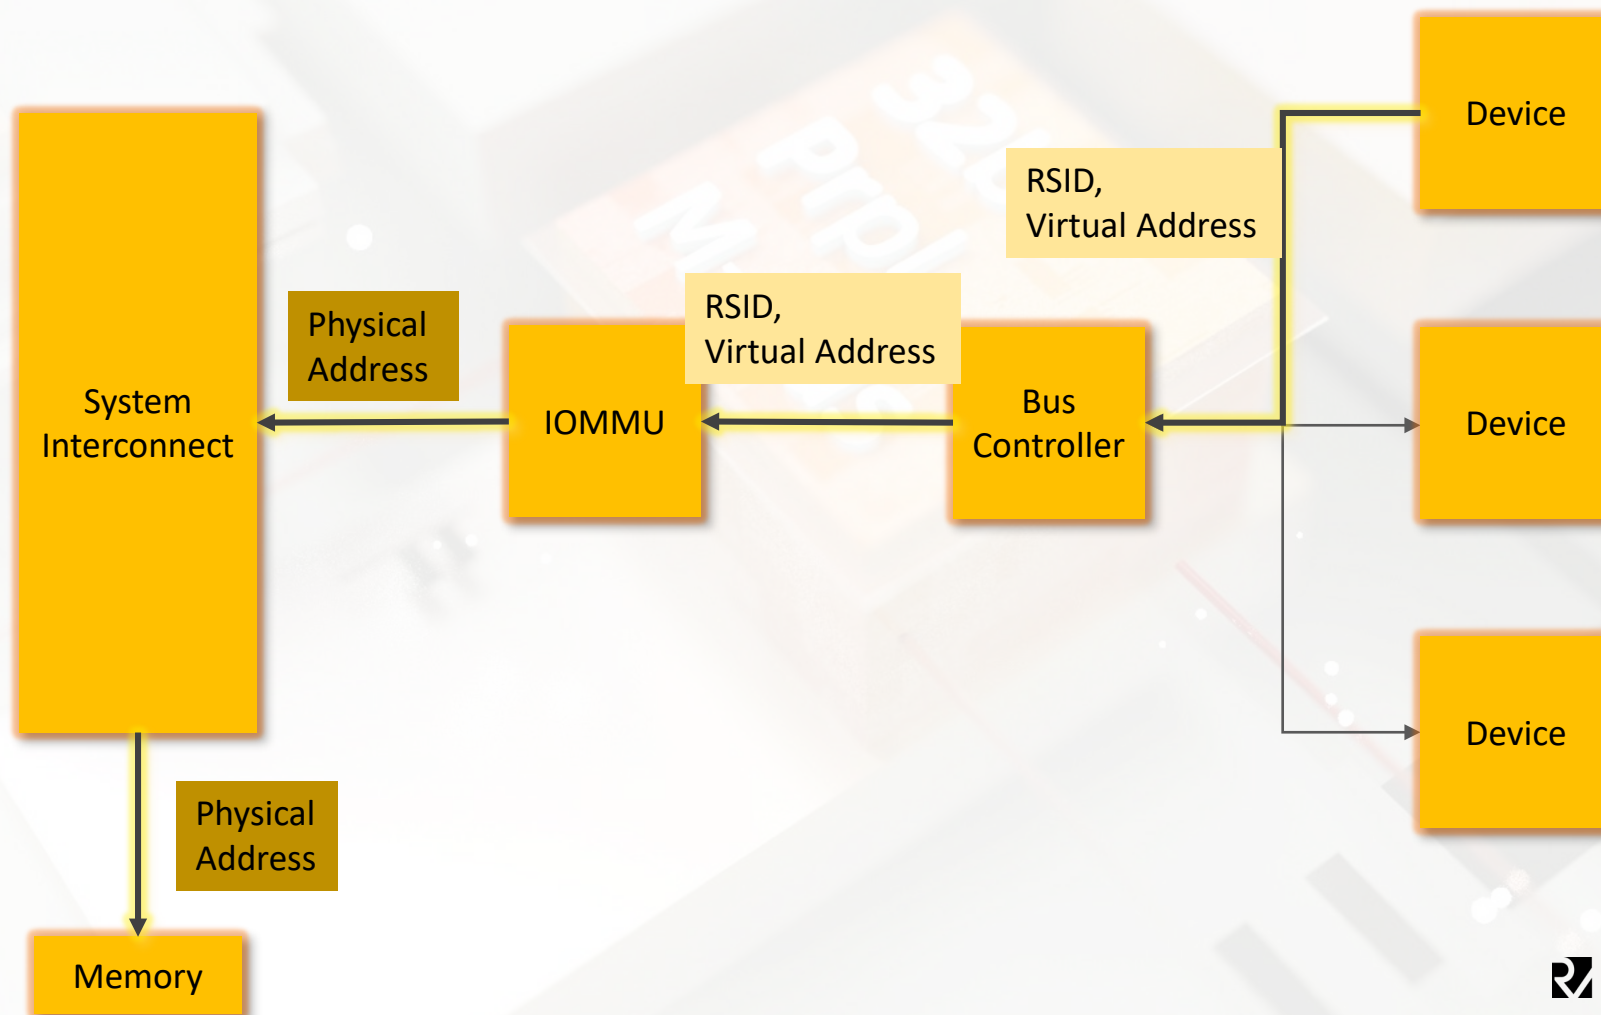
- Overview
- Features
- Prototype
- Future Work

Overview

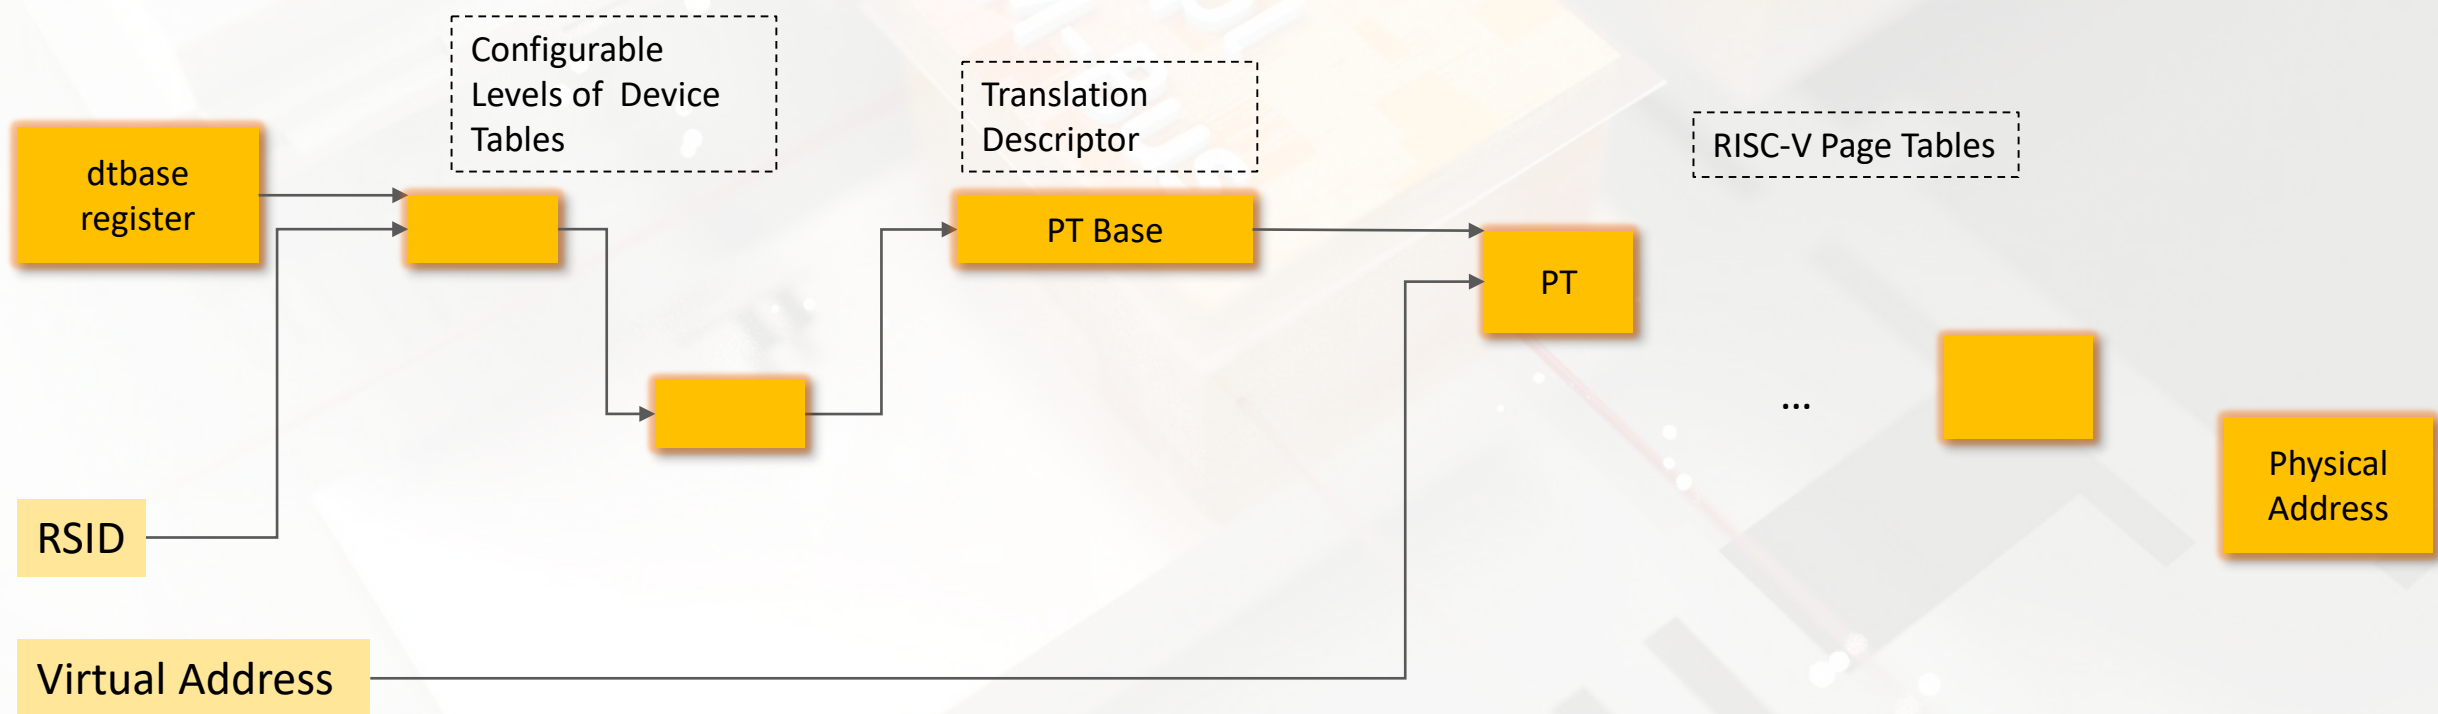
- IOMMU is needed for RISC-V
 - Kernel protection
 - User space drivers
 - Virtualization
 - Security
- T-Head release a draft of its IOMMU spec in 2021, covering essential features for an IOMMU for RISC-V
 - Entered as one of the proposals for the IOMMU TG's spec
 - Link: <https://github.com/sqzsq/xuantie-iommu-spec>
- An early prototype on QEMU for OS protection and user space driver was release in early 2022
- Since then, we have validated our design on a device passthrough prototype, and introduced additional features

Features

- An IOMMU translates the addresses in the DMA requests from the peripheral devices

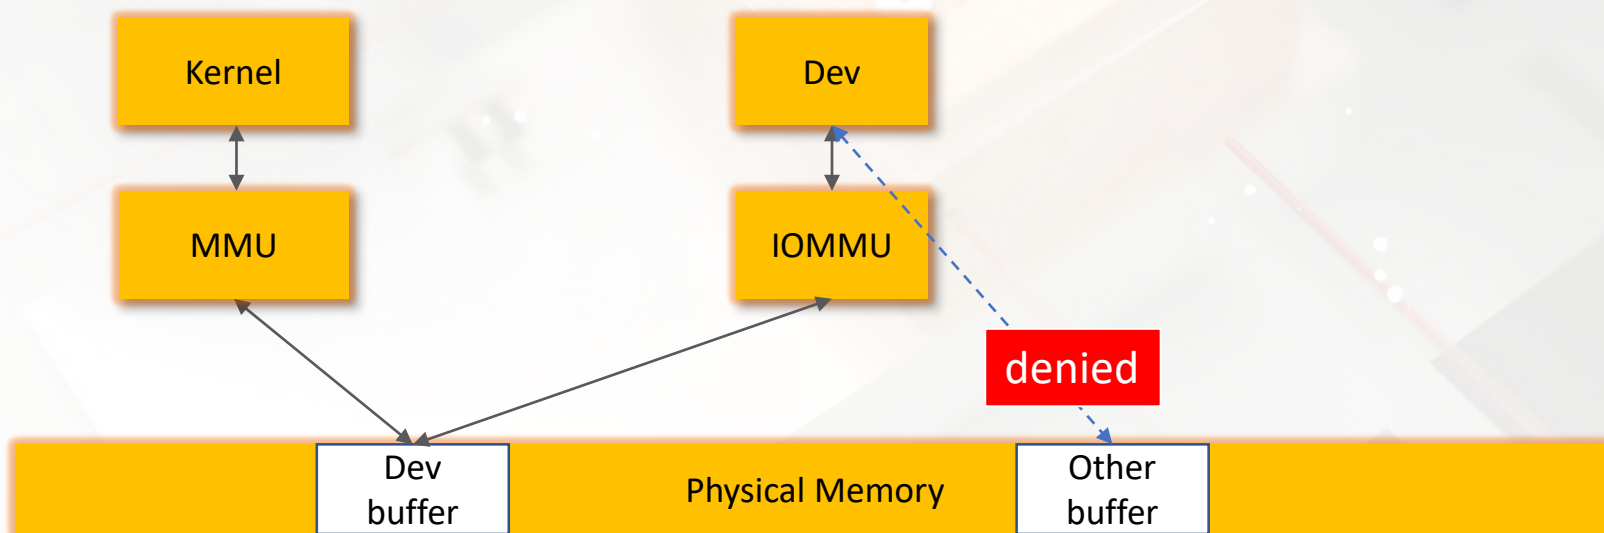


- The device table can be easily scaled up by adding more levels
 - Depending on how the devices under the management of the IOMMU is identified
- Separate the device context from the device tables for flexibility
 - Provides possibility for software to separately manage the two
- Compatible with RISC-V page tables for address translation

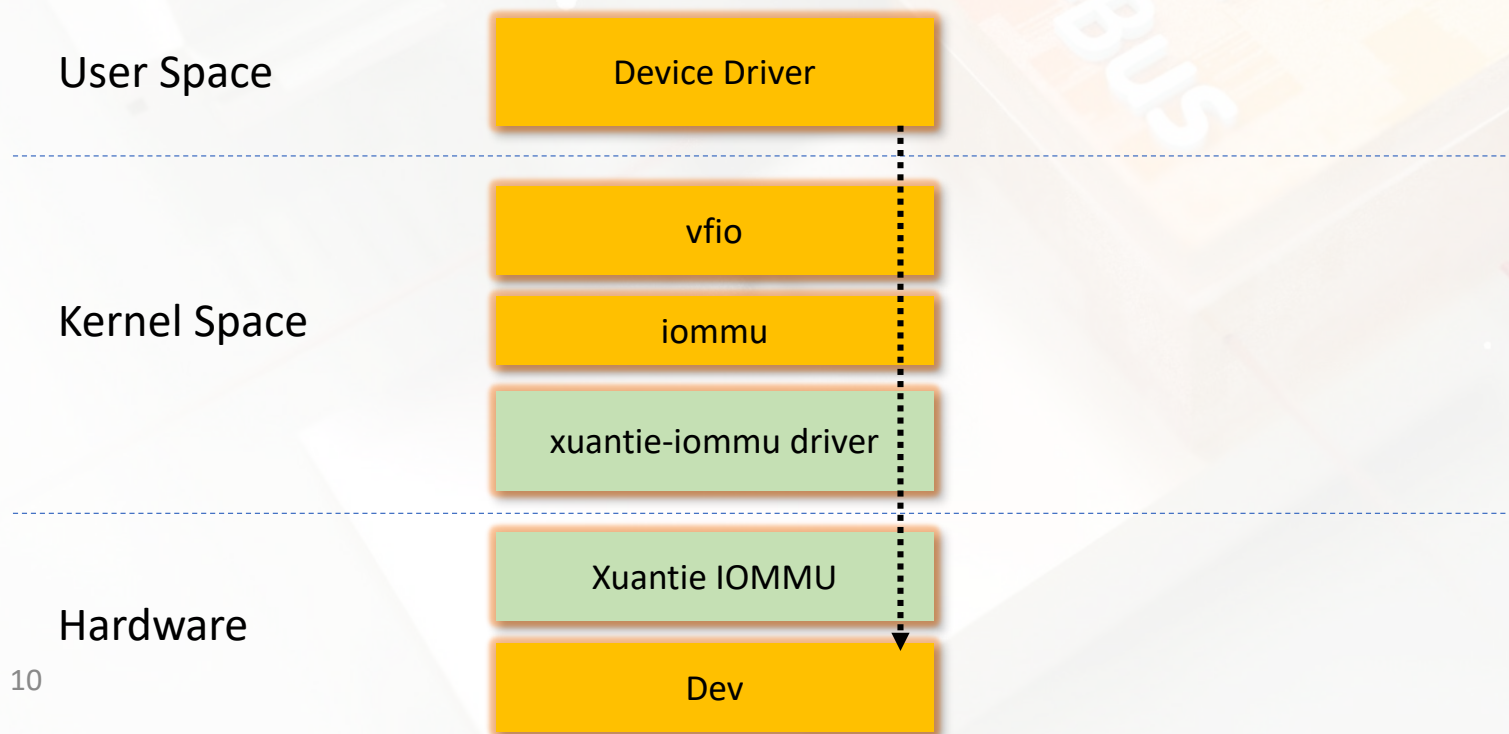


- There are multiple ways to implement the IOMMU in hardware
 - Private to device
 - Monolithic and shared
 - Distributed and shared
 - ...
- One possible way to integrate the IOMMU is a distributed implementation
- Each device is equipped with a device-specific TLB
- A shared TLB sits with the Page Table Walk logic, serving as a second level of TLB

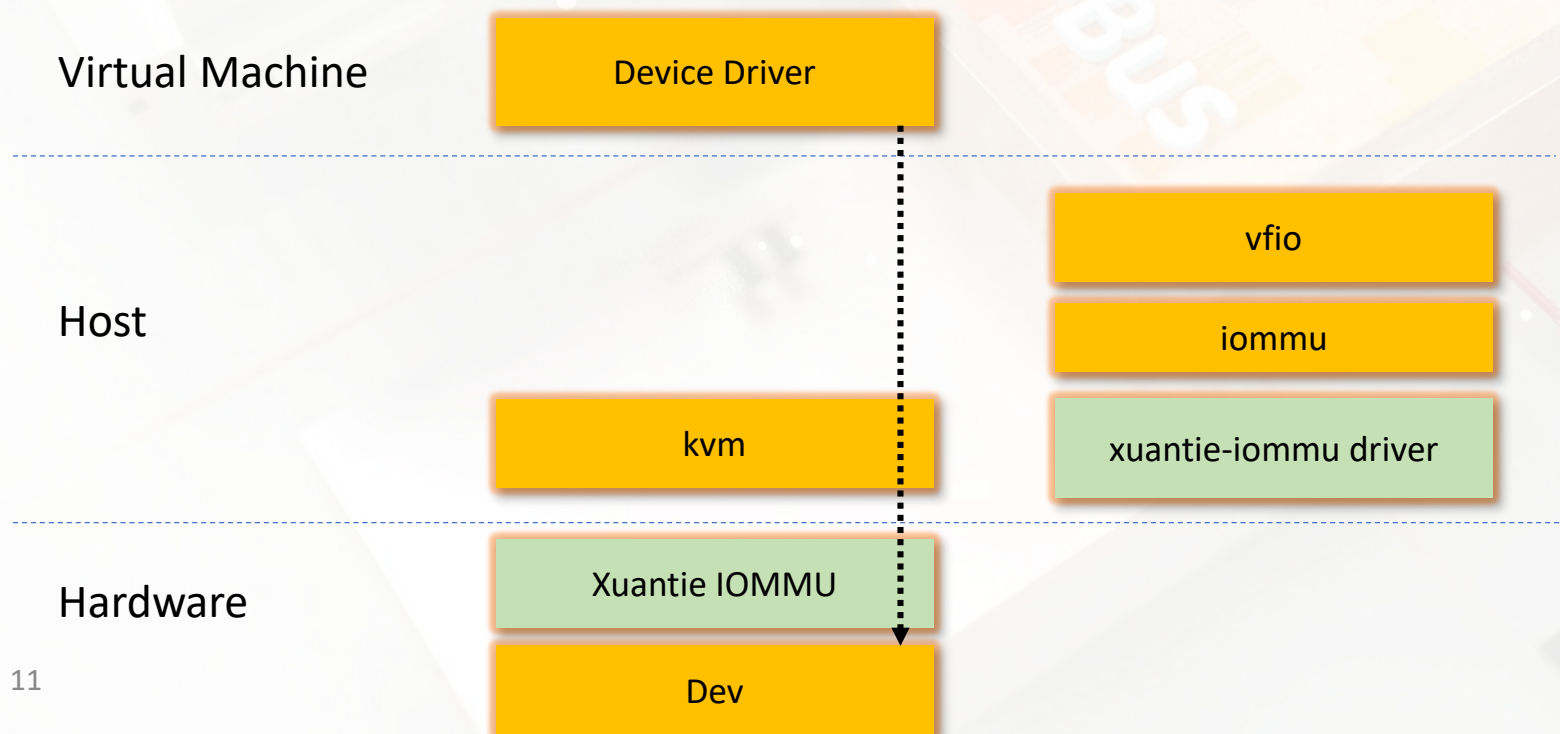
- The OS kernel utilize the IOMMU to impose restriction on the range of physical address a device can access
 - Security
 - Reliability
 - Debug
 - Memory utilization
 - ...
- The Linux kernel's IOMMU layer provides APIs for kernel space drivers to manage the buffer
- Our IOMMU driver fully conforms to the Linux kernel IOMMU layer, no change to the device driver is required.



- The OS kernel can also let a user space application directly control a device, in which case a user space driver is execution, for example, DPDK and SPDK etc.
- The existing VFIO framework allows user space programs to directly operation on PCIe devices and platform devices.
- Our IOMMU prototype is VFIO enabled and runs a user space driver
- No other change is required for the user space driver

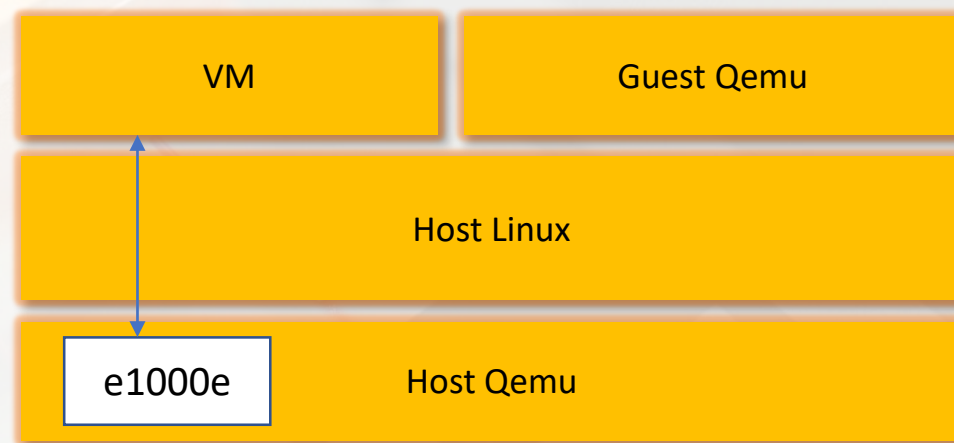


- For efficiency, virtual machines is given direct access to peripheral devices, i.e. passthrough devices.
- DMA from these devices are restricted to the memory of the virtual machine it belongs to.
 - IOMMU's interposing on the DMA requests can achieve this.
- Our IOMMU fits into the KVM/VFIO facility for device passthrough
- No change in the KVM module or VFIO framework is required



Prototype

- We have evaluated our design on QEMU + Linux for the following use cases
 - Host usage
 - Device passthrough
- For device passthrough we enabled the hypervisor extension on QEMU and setup KVM
 - The passthrough device is an e1000e NIC emulated by host-QEMU
 - We enabled necessary VFIO facility to let the VM directly use the NIC



- Host usage demo

- device passthrough demo

Future Work

- More features
 - PCIe support
 - AIA integration
 - Debug / trace
- Virtual IOMMU
- Hardware implementation
- Security / Confidential Computing
- Device co-operation

END